

# Лабораторная работа №1

## Логистическая регрессия в качестве нейронной сети

В работе предлагается использовать набор данных *notMNIST*, который состоит из изображений размерностью  $28 \times 28$  первых 10 букв латинского алфавита (*A* ... *J*, соответственно). Обучающая выборка содержит порядка 500 тыс. изображений, а тестовая – около 19 тыс.

Данные можно скачать по ссылке:

- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz)) (большой набор данных);
- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz)) (маленький набор данных);

Описание данных на английском языке доступно по ссылке: <http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html> (<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>).

### Задание 1

Загрузите данные и отобразите на экране несколько из изображений с помощью языка Python.

In [0]:

```
SMALL_DS_URL = 'https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz'  
LARGE_DS_URL = 'https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz'
```

In [0]:

```
%matplotlib inline  
  
import matplotlib.pyplot as plt
```

In [0]:

```
from urllib.request import urlretrieve
import tarfile
import os

def tar_to_dir(_tar_url, _key):
    dir_name_ = 'dataset_' + _key
    local_file_name_ = dir_name_ + '.f'

    urlretrieve(_tar_url, local_file_name_)

    with tarfile.open(local_file_name_, 'r:gz') as tar_:
        tar_.extractall(dir_name_)

    os.remove(local_file_name_)

    return dir_name_
```

In [0]:

```
def get_examples(_dataframe, _label_column_name, _data_column_name):
    n_ = _dataframe[_label_column_name].nunique()

    examples_ = _dataframe.sample(n_)[_data_column_name]

    return examples_
```

In [0]:

```
from math import ceil
import numpy as np

def print_examples(_examples):
    fig = plt.figure(figsize = (16, 6))

    height_ = 2
    width_ = ceil(_examples.count() / height_)

    for i, item_ in enumerate(_examples):

        ax = fig.add_subplot(height_, width_, i + 1)
        ax.axis('off')
        ax.imshow(item_, cmap = 'gray', interpolation = 'none')

    plt.show()
```

In [0]:

```
from imageio import imread
import pandas as pd

def image_to_array(_image):
    try:
        array_ = imread(_image)

        return True, array_
    except:
        return False, None

def get_inner_dir(_dir_path):
    return [x[0] for x in os.walk(_dir_path)][1]

def remove_duplicates(_dataframe, _data_column_name):
    return _dataframe.loc[_dataframe[_data_column_name].astype(str).drop_duplicates().index]

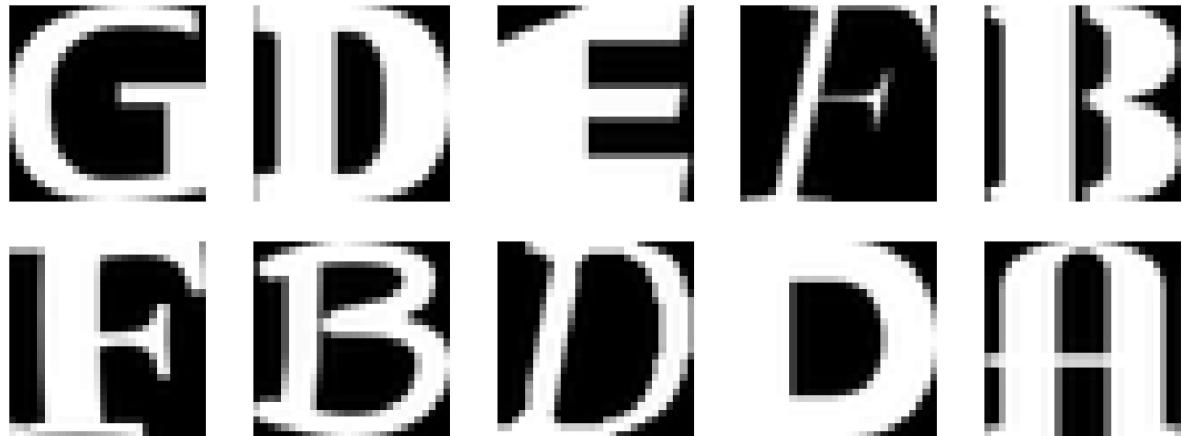
def dir_to_dataframe(_dir_path):
    dataframes_ = []
    inner_dir_path_ = get_inner_dir(_dir_path)
    for subdir_ in sorted(os.listdir(inner_dir_path_)):
        letter_ = subdir_
        data_ = []
        files_ = os.listdir(os.path.join(inner_dir_path_, subdir_))
        for f in files_:
            file_path_ = os.path.join(inner_dir_path_, subdir_, f)
            can_read_, im = image_to_array(file_path_)
            if can_read_:
                data_.append(im)
        g = [letter_] * len(data_)
        e = np.array(data_)
        h = pd.DataFrame()
        h['data'] = data_
        h['label'] = letter_
        dataframes_.append(h)
    result_ = pd.concat(dataframes_, ignore_index = True)
    unique_ = remove_duplicates(result_, 'data')
    return unique_
```

In [0]:

```
def tar_to_dataframe(_tar_url, _key):  
    dir_name_ = tar_to_dir(_tar_url, _key)  
    inner_dir_ = get_inner_dir(dir_name_)  
    dataframe_ = dir_to_dataframe(dir_name_)  
    examples_ = get_examples(dataframe_, 'label', 'data')  
    print_examples(examples_)  
  
    return dataframe_
```

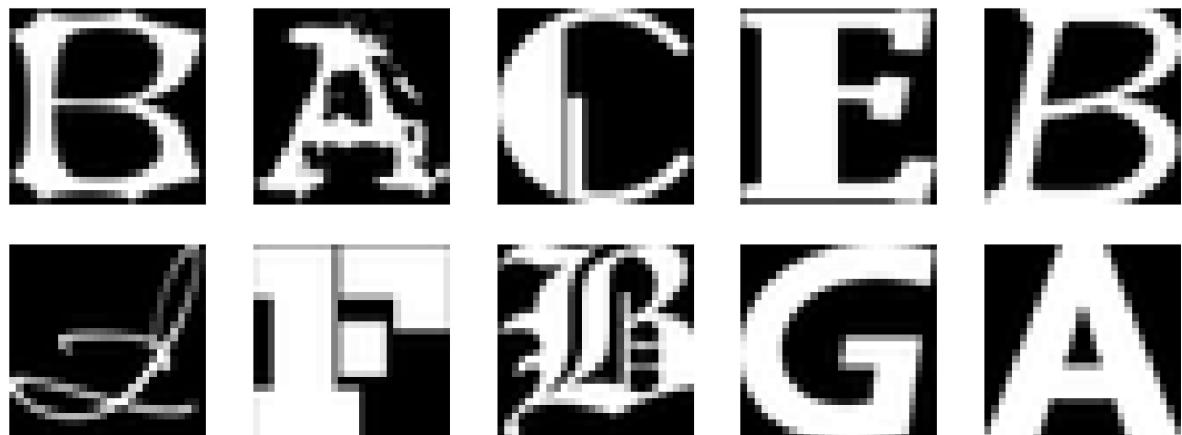
In [8]:

```
small_dataframe = tar_to_dataframe(SMALL_DS_URL, 'small')
```



In [9]:

```
large_dataframe = tar_to_dataframe(LARGE_DS_URL, 'large')
```



## Задание 2

Проверьте, что классы являются сбалансированными, т.е. количество изображений, принадлежащих каждому из классов, примерно одинаково (в данной задаче 10 классов).

In [0]:

```
def print_balance(_dataframe, _label_column_name):  
    values_ = _dataframe[_label_column_name].value_counts().sort_values(ascending = False)  
    print((':>10') * len(values_)).format(*values_))
```

In [11]:

```
print_balance(small_dataframe, 'label')
```

1853	1850	1850	1848	1848	1848	1847	1847
1847	1845	1596					

In [12]:

```
print_balance(large_dataframe, 'label')
```

47226	47102	47012	46890	46771	46663	46577	4
6521	46098	41086					

Как видим, классы сбалансированы.

### Задание 3

Разделите данные на три подвыборки: обучающую (200 тыс. изображений), валидационную (10 тыс. изображений) и контрольную (тестовую) (19 тыс. изображений).

In [0]:

```
def split(_dataframe, _n_train, _n_test, _n_val):  
  
    assert _dataframe.shape[0] >= _n_train + _n_test + _n_val  
  
    to_be_split_ = _dataframe.copy(deep = True)  
  
    seed_ = 666  
  
    train_ = to_be_split_.sample(n = _n_train, random_state = seed_)  
  
    to_be_split_ = to_be_split_.drop(train_.index)  
    test_ = to_be_split_.sample(n = _n_test, random_state = seed_)  
  
    val_ = to_be_split_.drop(test_.index).sample(n = _n_val, random_state = seed_)  
  
    return train_, test_, val_
```

In [14]:

```
large_dataframe.shape[0]
```

Out[14]:

461946

In [15]:

```
train, test, validation = split(large_dataframe, 200000, 10000, 19000)

print_balance(train, 'label')
print_balance(test, 'label')
print_balance(validation, 'label')
```

	20415	20350	20317	20290	20278	20191	20170	2
0124	20100	17765						
	1049	1043	1033	1029	1021	1016	995	
981	961	872						
	2014	1945	1934	1931	1912	1903	1898	
1887	1864	1712						

Видно, что удалось сохранить баланс между классами.

## Задание 4

Проверьте, что данные из обучающей выборки не пересекаются с данными из валидационной и контрольной выборок. Другими словами, избавьтесь от дубликатов в обучающей выборке.

In [0]:

```
def no_duplicates(_dataframe, _data_column_name):

    original_length_ = _dataframe.shape[0]

    unique_length_ = _dataframe[_data_column_name].astype(str).unique().shape[0]

    print(str(original_length_) + ' -- ' + str(unique_length_))

    return original_length_ == unique_length_
```

In [17]:

```
print(no_duplicates(small_dataframe, 'data'))
```

```
18232 -- 18232
True
```

In [18]:

```
print(no_duplicates(large_dataframe, 'data'))
```

```
461946 -- 461946
True
```

In [0]:

```
small_dataframe.to_pickle("./small.pkl")
large_dataframe.to_pickle("./large.pkl")
```

Дубликатов не обнаружено, так как они были удалены на шаге построения датасета из файлов.

## Задание 5

Постройте простейший классификатор (например, с помощью логистической регрессии). Постройте график зависимости точности классификатора от размера обучающей выборки (50, 100, 1000, 50000). Для построения классификатора можете использовать библиотеку *SkLearn* (<http://scikit-learn.org>). (<http://scikit-learn.org>).

In [0]:

```
def dataframe_to_x_y(_dataframe):  
    x_ = np.stack(_dataframe['data']).reshape((_dataframe.shape[0], -1))  
    y_ = _dataframe['label'].to_numpy()  
  
    return x_, y_
```

In [0]:

```
X_train, y_train = dataframe_to_x_y(train)  
X_test, y_test = dataframe_to_x_y(test)
```

In [0]:

```
sizes = [50, 100, 1000, 50000]  
  
clfs = {}  
  
scores = {}
```

In [23]:

```
from sklearn.linear_model import LogisticRegression

for size_ in sizes:

    clf_ = LogisticRegression(max_iter = 100).fit(X_train[:size_], y_train[:size_])

    clfs[size_] = clf_
```

/usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:94  
0: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

/usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:94  
0: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

/usr/local/lib/python3.6/dist-packages/sklearn/linear\_model/\_logistic.py:94  
0: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

In [24]:

```
print(*clfs[50000].predict(X_test[:10]), sep = '\t')
```

E C B H C G E C G H

In [25]:

```
print(*y_test[:10], sep = '\t')
```

C C B H C G B E G H

In [0]:

```
for size_ in sizes:  
    scores[size_] = clfs[size_].score(X_test, y_test)
```

In [27]:

```
print(scores)
```

```
{50: 0.559, 100: 0.6609, 1000: 0.7231, 50000: 0.8155}
```

In [28]:

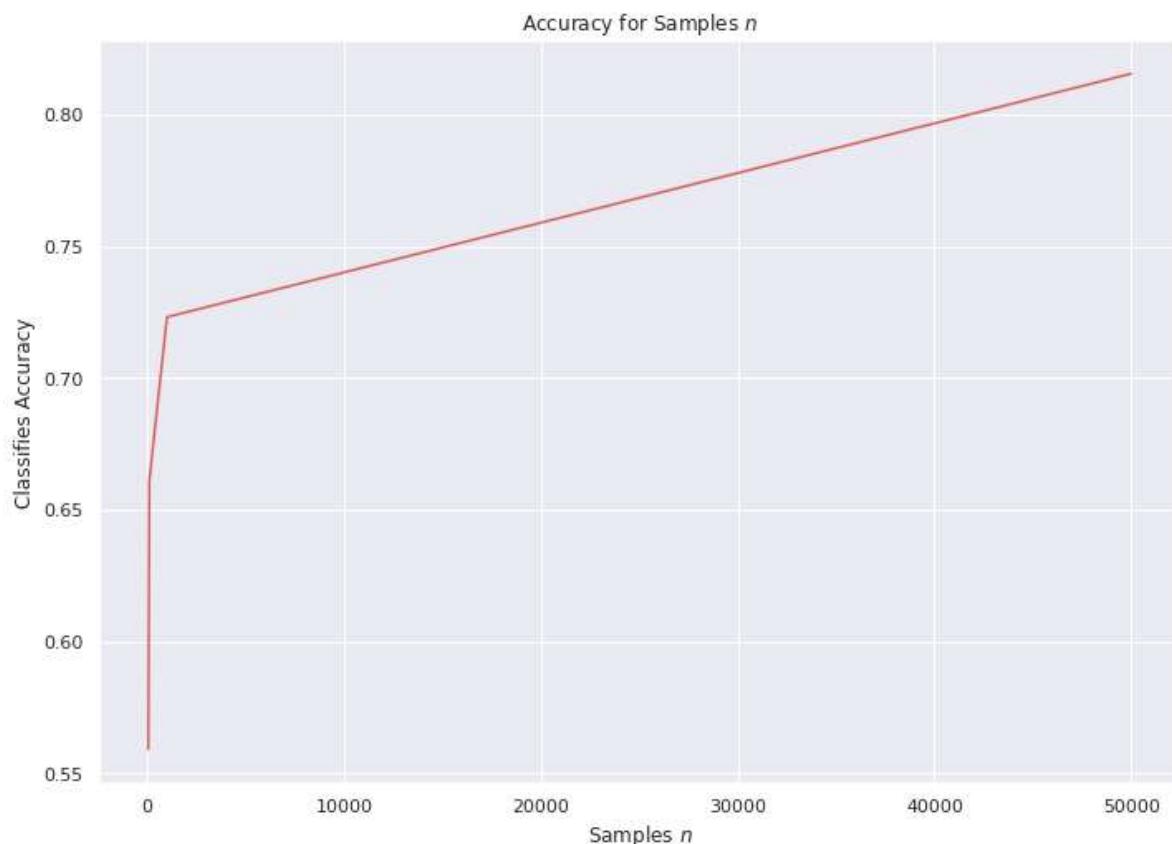
```
import seaborn as sns  
  
from matplotlib import rcParams  
  
rcParams['figure.figsize'] = 11.7, 8.27  
  
sns.set()  
  
sns.set_palette(sns.color_palette('hls'))
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
```

```
    import pandas.util.testing as tm
```

In [29]:

```
sns.lineplot(sizes, [scores[s] for s in sizes])  
  
plt.xlabel('Samples $n$')  
plt.ylabel('Classifies Accuracy')  
  
plt.title('Accuracy for Samples $n$')  
  
plt.show()
```



На графике видим, что с увеличением выборки качество классификации растёт как логарифмическая функция от размера выборки.

# Лабораторная работа №2

## Реализация глубокой нейронной сети

В работе предлагается использовать набор данных *notMNIST*, который состоит из изображений размерностью  $28 \times 28$  первых 10 букв латинского алфавита (*A* ... *J*, соответственно). Обучающая выборка содержит порядка 500 тыс. изображений, а тестовая – около 19 тыс.

Данные можно скачать по ссылке:

- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz)) (большой набор данных);
- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz)) (маленький набор данных);

Описание данных на английском языке доступно по ссылке: <http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html> (<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>).

### Задание 1

Реализуйте полносвязную нейронную сеть с помощью библиотеки *TensorFlow*. В качестве алгоритма оптимизации можно использовать, например, стохастический градиент (*Stochastic Gradient Descent*, *SGD*). Определите количество скрытых слоев от 1 до 5, количество нейронов в каждом из слоев до нескольких сотен, а также их функции активации (кусочно-линейная, сигмоидная, гиперболический тангенс и т.д.).

In [1]:

```
from google.colab import drive  
  
drive.mount('/content/drive', force_remount = True)
```

Go to this URL in a browser: [\(https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response\\_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly\)](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly (https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly))

Enter your authorization code:

.....

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2'

import sys

sys.path.append(BASE_DIR)

import os

os.chdir(BASE_DIR)
```

In [0]:

```
import pandas as pd

dataframe = pd.read_pickle("./large.pkl")
```

In [4]:

```
dataframe['data'].shape
```

Out[4]:

```
(461946,)
```

In [5]:

```
! pip install tensorflow-gpu --pre --quiet

! pip show tensorflow-gpu
```

```
|██████████| 516.2MB 26kB/s
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: termcolor, tensorflow-estimator, opt-einsum, wrapt, absl-py, protobuf, tensorboard, astunparse, keras-preprocessing, h5py, numpy, wheel, scipy, gast, six, google-pasta, grpcio
Required-by:
```

In [0]:

```
import tensorflow as tf
```

In [0]:

```
import numpy as np
```

In [0]:

```
dataframe_test = dataframe.sample(frac = 0.1)

dataframe = dataframe.drop(dataframe_test.index)
```

In [9]:

```
x = np.asarray(list(dataframe['data']))[..., np.newaxis]

x = tf.keras.utils.normalize(x, axis = 1)

x.shape
```

Out[9]:

```
(415751, 28, 28, 1)
```

In [31]:

```
x_test = np.asarray(list(dataframe_test['data']))[..., np.newaxis]

x_test = tf.keras.utils.normalize(x_test, axis = 1)

x_test.shape
```

Out[31]:

```
(46195, 28, 28, 1)
```

In [0]:

```
IMAGE_DIM_0, IMAGE_DIM_1 = x.shape[1], x.shape[2]
```

In [11]:

```
from tensorflow.keras.utils import to_categorical

y = to_categorical(dataframe['label'].astype('category').cat.codes.astype('int32'))

y.shape
```

Out[11]:

```
(415751, 10)
```

In [32]:

```
y_test = to_categorical(dataframe_test['label'].astype('category').cat.codes.astype('int32')

y_test.shape
```

Out[32]:

```
(46195, 10)
```

In [0]:

```
LAYER_WIDTH = 5000
```

In [0]:

```
CLASSES_N = y.shape[1]
```

In [0]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Reshape

model = tf.keras.Sequential()

model.add(Reshape((IMAGE_DIM_0 * IMAGE_DIM_1,), input_shape = (IMAGE_DIM_0, IMAGE_DIM_1, 1)))
model.add(Dense(LAYER_WIDTH, activation = 'relu'))
model.add(Dense(LAYER_WIDTH, activation = 'sigmoid'))
model.add(Dense(LAYER_WIDTH, activation = 'tanh'))
model.add(Dense(LAYER_WIDTH, activation = 'elu'))
model.add(Dense(LAYER_WIDTH, activation = 'softmax'))
model.add(Dense(CLASSES_N))
```

In [0]:

```
def cat_cross_from_logits(y_true, y_pred):
    return tf.keras.losses.categorical_crossentropy(y_true, y_pred, from_logits = True)

model.compile(optimizer = 'sgd',
              loss = cat_cross_from_logits,
              metrics = ['categorical_accuracy'])
```

In [16]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
reshape (Reshape)	(None, 784)	0
dense (Dense)	(None, 5000)	3925000
dense_1 (Dense)	(None, 5000)	25005000
dense_2 (Dense)	(None, 5000)	25005000
dense_3 (Dense)	(None, 5000)	25005000
dense_4 (Dense)	(None, 5000)	25005000
dense_5 (Dense)	(None, 10)	50010
<hr/>		

Total params: 103,995,010

Trainable params: 103,995,010

Non-trainable params: 0

In [0]:

```
VAL_SPLIT_RATE = 0.1
```

In [0]:

```
EPOCHS_N = 10
```

In [19]:

```
model.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)
```

```
Epoch 1/10
11693/11693 [=====] - 103s 9ms/step - loss: 2.2195
- categorical_accuracy: 0.1122 - val_loss: 4.8722 - val_categorical_accuracy: 0.0000e+00
Epoch 2/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2042
- categorical_accuracy: 0.1130 - val_loss: 5.4886 - val_categorical_accuracy: 0.0000e+00
Epoch 3/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2025
- categorical_accuracy: 0.1126 - val_loss: 5.8407 - val_categorical_accuracy: 0.0000e+00
Epoch 4/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2018
- categorical_accuracy: 0.1127 - val_loss: 6.0791 - val_categorical_accuracy: 0.0000e+00
Epoch 5/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2015
- categorical_accuracy: 0.1123 - val_loss: 6.2565 - val_categorical_accuracy: 0.0000e+00
Epoch 6/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2013
- categorical_accuracy: 0.1127 - val_loss: 6.3940 - val_categorical_accuracy: 0.0000e+00
Epoch 7/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2012
- categorical_accuracy: 0.1134 - val_loss: 6.5053 - val_categorical_accuracy: 0.0000e+00
Epoch 8/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2011
- categorical_accuracy: 0.1130 - val_loss: 6.5967 - val_categorical_accuracy: 0.0000e+00
Epoch 9/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2010
- categorical_accuracy: 0.1134 - val_loss: 6.6732 - val_categorical_accuracy: 0.0000e+00
Epoch 10/10
11693/11693 [=====] - 102s 9ms/step - loss: 2.2010
- categorical_accuracy: 0.1134 - val_loss: 6.7377 - val_categorical_accuracy: 0.0000e+00
```

Out[19]:

```
<tensorflow.python.keras.callbacks.History at 0x7f399433fbe0>
```

In [34]:

```
results = model.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results)
```

```
1444/1444 [=====] - 5s 4ms/step - loss: 2.6571 - categorical_accuracy: 0.1008
Test loss, test accuracy: [2.657094955444336, 0.10081177949905396]
```

## Задание 2

Как улучшилась точность классификатора по сравнению с логистической регрессией?

Стало хуже — на тестовой выборке точность составила 10%. Похоже, что данная модель совершенно не подходит для решения этой задачи.

## Задание 3

Используйте регуляризацию и метод сброса нейронов (*dropout*) для борьбы с переобучением. Как улучшилось качество классификации?

In [0]:

```
REG_RATE = 0.001
```

In [0]:

```
from tensorflow.keras.regularizers import l2
l2_reg = l2(REG_RATE)
```

In [0]:

```
DROPOUT_RATE = 0.2
```

In [0]:

```
from tensorflow.keras.layers import Dropout
dropout_layer = Dropout(DROPOUT_RATE)
```

In [0]:

```
model_2 = tf.keras.Sequential()

model_2.add(Reshape((IMAGE_DIM_0 * IMAGE_DIM_1,), input_shape = (IMAGE_DIM_0, IMAGE_DIM_1,
model_2.add(Dense(LAYER_WIDTH, activation = 'relu', kernel_regularizer = l2_reg))
model_2.add(dropout_layer)
model_2.add(Dense(LAYER_WIDTH, activation = 'sigmoid', kernel_regularizer = l2_reg))
model_2.add(dropout_layer)
model_2.add(Dense(LAYER_WIDTH, activation = 'tanh', kernel_regularizer = l2_reg))
model_2.add(dropout_layer)
model_2.add(Dense(LAYER_WIDTH, activation = 'sigmoid', kernel_regularizer = l2_reg))
model_2.add(dropout_layer)
model_2.add(Dense(LAYER_WIDTH, activation = 'relu', kernel_regularizer = l2_reg))
model_2.add(dropout_layer)
model_2.add(Dense(CLASSES_N))
```

In [0]:

```
model_2.compile(optimizer = 'sgd',
                 loss = cat_crossentropy,
                 metrics = ['categorical_accuracy'])
```

In [26]:

```
model_2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
reshape_1 (Reshape)	(None, 784)	0
dense_6 (Dense)	(None, 5000)	3925000
dropout (Dropout)	(None, 5000)	0
dense_7 (Dense)	(None, 5000)	25005000
dense_8 (Dense)	(None, 5000)	25005000
dense_9 (Dense)	(None, 5000)	25005000
dense_10 (Dense)	(None, 5000)	25005000
dense_11 (Dense)	(None, 10)	50010
=====		
Total params: 103,995,010		
Trainable params: 103,995,010		
Non-trainable params: 0		

In [27]:

```
model_2.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)

Epoch 1/10
11693/11693 [=====] - 210s 18ms/step - loss: 19.270
- categorical_accuracy: 0.1154 - val_loss: 19.6030 - val_categorical_accuracy: 0.0000e+00
Epoch 2/10
11693/11693 [=====] - 210s 18ms/step - loss: 12.889
- categorical_accuracy: 0.1193 - val_loss: 15.0547 - val_categorical_accuracy: 0.0000e+00
Epoch 3/10
11693/11693 [=====] - 210s 18ms/step - loss: 8.4027
- categorical_accuracy: 0.3462 - val_loss: 12.9659 - val_categorical_accuracy: 0.0000e+00
Epoch 4/10
11693/11693 [=====] - 210s 18ms/step - loss: 5.6802
- categorical_accuracy: 0.4805 - val_loss: 11.1155 - val_categorical_accuracy: 0.0000e+00
Epoch 5/10
11693/11693 [=====] - 210s 18ms/step - loss: 4.0491
- categorical_accuracy: 0.5286 - val_loss: 9.5503 - val_categorical_accuracy: 0.0000e+00
Epoch 6/10
11693/11693 [=====] - 210s 18ms/step - loss: 3.0297
- categorical_accuracy: 0.5674 - val_loss: 8.9475 - val_categorical_accuracy: 0.0000e+00
Epoch 7/10
11693/11693 [=====] - 210s 18ms/step - loss: 2.3880
- categorical_accuracy: 0.5944 - val_loss: 8.4980 - val_categorical_accuracy: 0.0000e+00
Epoch 8/10
11693/11693 [=====] - 210s 18ms/step - loss: 1.9772
- categorical_accuracy: 0.6292 - val_loss: 7.8002 - val_categorical_accuracy: 0.0000e+00
Epoch 9/10
11693/11693 [=====] - 210s 18ms/step - loss: 1.7196
- categorical_accuracy: 0.6475 - val_loss: 7.6441 - val_categorical_accuracy: 0.0000e+00
Epoch 10/10
11693/11693 [=====] - 211s 18ms/step - loss: 1.5578
- categorical_accuracy: 0.6599 - val_loss: 7.5225 - val_categorical_accuracy: 0.0000e+00
```

Out[27]:

```
<tensorflow.python.keras.callbacks.History at 0x7f399345a9b0>
```

In [36]:

```
results_2 = model_2.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results_2)
```

```
1444/1444 [=====] - 8s 6ms/step - loss: 1.4656 - categorical_accuracy: 0.7228
Test loss, test accuracy: [1.4655554294586182, 0.7227838635444641]
```

Регуляризация и сброс нейронов значительно помогли — модель показывает 72% точности на тестовой

выборке!

## Задание 4

Воспользуйтесь динамически изменяемой скоростью обучения (*learning rate*). Наилучшая точность, достигнутая с помощью данной модели составляет 97.1%. Какую точность демонстрирует Ваша реализованная модель?

In [28]:

```
from tensorflow.keras.optimizers import SGD

dyn_lr_sgd = SGD(lr = 0.01, momentum = 0.9)

model_2.compile(optimizer = dyn_lr_sgd,
                 loss = cat_crossentropy,
                 metrics = ['categorical_accuracy'])

model_2.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)
```

```
Epoch 1/10
11693/11693 [=====] - 240s 20ms/step - loss: 1.3780
- categorical_accuracy: 0.6521 - val_loss: 6.3145 - val_categorical_accuracy: 0.0000e+00
Epoch 2/10
11693/11693 [=====] - 239s 20ms/step - loss: 1.0968
- categorical_accuracy: 0.7462 - val_loss: 6.1084 - val_categorical_accuracy: 0.0000e+00
Epoch 3/10
11693/11693 [=====] - 239s 20ms/step - loss: 1.0805
- categorical_accuracy: 0.7522 - val_loss: 6.2995 - val_categorical_accuracy: 0.0000e+00
Epoch 4/10
11693/11693 [=====] - 238s 20ms/step - loss: 1.0759
- categorical_accuracy: 0.7534 - val_loss: 6.1664 - val_categorical_accuracy: 0.0000e+00
Epoch 5/10
11693/11693 [=====] - 239s 20ms/step - loss: 1.0707
- categorical_accuracy: 0.7564 - val_loss: 6.4368 - val_categorical_accuracy: 0.0000e+00
Epoch 6/10
11693/11693 [=====] - 238s 20ms/step - loss: 1.0186
- categorical_accuracy: 0.7755 - val_loss: 6.6284 - val_categorical_accuracy: 0.0000e+00
Epoch 7/10
11693/11693 [=====] - 239s 20ms/step - loss: 0.9949
- categorical_accuracy: 0.7805 - val_loss: 6.6951 - val_categorical_accuracy: 0.0000e+00
Epoch 8/10
11693/11693 [=====] - 239s 20ms/step - loss: 0.9902
- categorical_accuracy: 0.7824 - val_loss: 6.1755 - val_categorical_accuracy: 0.0000e+00
Epoch 9/10
11693/11693 [=====] - 239s 20ms/step - loss: 0.9873
- categorical_accuracy: 0.7840 - val_loss: 6.3812 - val_categorical_accuracy: 0.0000e+00
Epoch 10/10
11693/11693 [=====] - 239s 20ms/step - loss: 0.9844
- categorical_accuracy: 0.7852 - val_loss: 6.3063 - val_categorical_accuracy: 0.0000e+00
```

Out[28]:

```
<tensorflow.python.keras.callbacks.History at 0x7f39922ee0f0>
```

In [38]:

```
results_3 = model_2.evaluate(x_test, y_test)
print('Test loss, test accuracy:', results_3)
```

```
1444/1444 [=====] - 8s 6ms/step - loss: 1.4656 - categorical_accuracy: 0.7228
Test loss, test accuracy: [1.465554294586182, 0.7227838635444641]
```

Динамически изменяемая скорость обучения не улучшила результат — 72% на тестовой выборке.

Можно сделать вывод, что модель с полносвязными слоями может использоваться для решения задачи распознавания изображений, однако она очевидно не является наилучшей.

# Лабораторная работа №3

## Реализация сверточной нейронной сети

В работе предлагается использовать набор данных *notMNIST*, который состоит из изображений размерностью  $28 \times 28$  первых 10 букв латинского алфавита (*A* ... *J*, соответственно). Обучающая выборка содержит порядка 500 тыс. изображений, а тестовая – около 19 тыс.

Данные можно скачать по ссылке:

- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_large.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_large.tar.gz)) (большой набор данных);
- [https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz) ([https://commondatastorage.googleapis.com/books1000/notMNIST\\_small.tar.gz](https://commondatastorage.googleapis.com/books1000/notMNIST_small.tar.gz)) (маленький набор данных);

Описание данных на английском языке доступно по ссылке: <http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html> (<http://yaroslavvb.blogspot.sg/2011/09/notmnist-dataset.html>).

### Задание 1

Реализуйте нейронную сеть с двумя сверточными слоями, и одним полно связанным с нейронами с кусочно-линейной функцией активации. Какова точность построенной модели?

In [1]:

```
from google.colab import drive  
drive.mount('/content/drive', force_remount = True)
```

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2'  
  
import sys  
  
sys.path.append(BASE_DIR)  
  
import os  
  
os.chdir(BASE_DIR)
```

In [0]:

```
import pandas as pd  
  
dataframe = pd.read_pickle("./large.pkl")
```

In [4]:

```
! pip install tensorflow-gpu --pre --quiet
```

```
! pip show tensorflow-gpu
```

```
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: astunparse, six, google-pasta, tensorflow-estimator, gast, tensorboard, grpcio, scipy, termcolor, wrapt, keras-preprocessing, protobuf, absl-py, numpy, h5py, opt-einsum, wheel
Required-by:
```

In [0]:

```
import tensorflow as tf
```

In [0]:

```
# To fix memory leak: https://github.com/tensorflow/tensorflow/issues/33009
```

```
tf.compat.v1.disable_eager_execution()
```

In [0]:

```
import numpy as np
```

In [0]:

```
dataframe_test = dataframe.sample(frac = 0.1)
dataframe = dataframe.drop(dataframe_test.index)
```

In [9]:

```
x = np.asarray(list(dataframe['data']))[..., np.newaxis]
x = tf.keras.utils.normalize(x, axis = 1)
x.shape
```

Out[9]:

```
(415751, 28, 28, 1)
```

In [10]:

```
x_test = np.asarray(list(dataframe_test['data']))[..., np.newaxis]  
x_test = tf.keras.utils.normalize(x_test, axis = 1)  
x_test.shape
```

Out[10]:

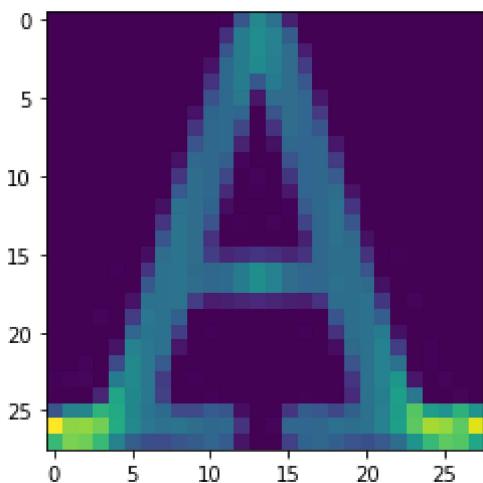
(46195, 28, 28, 1)

In [11]:

```
import matplotlib.pyplot as plt  
  
plt.imshow(x[100].squeeze())
```

Out[11]:

<matplotlib.image.AxesImage at 0x7f6ba8e81a90>



In [0]:

```
IMAGE_DIM_0, IMAGE_DIM_1 = x.shape[1], x.shape[2]
```

In [13]:

```
from tensorflow.keras.utils import to_categorical  
  
y = to_categorical(dataframe['label'].astype('category').cat.codes.astype('int32'))  
y.shape
```

Out[13]:

(415751, 10)

In [14]:

```
y_test = to_categorical(dataframe_test['label'].astype('category').cat.codes.astype('int32')
y_test.shape
```

Out[14]:

```
(46195, 10)
```

In [0]:

```
CLASSES_N = y.shape[1]
```

In [0]:

```
DENSE_LAYER_WIDTH = 5000
```

In [17]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Dense, Flatten

model = tf.keras.Sequential()

model.add(Conv2D(16, 3, padding = 'same', activation = 'relu', input_shape = (IMAGE_DIM_0,
model.add(Conv2D(32, 3, padding = 'same', activation = 'relu'))
model.add(Flatten())
model.add(Dense(DENSE_LAYER_WIDTH, activation = 'relu'))
model.add(Dense(CLASSES_N))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/resource\_variable\_ops.py:1666: calling BaseResourceVariable.\_\_init\_\_(from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

In [0]:

```
def cat_cross_from_logits(y_true, y_pred):
    return tf.keras.losses.categorical_crossentropy(y_true, y_pred, from_logits = True)

model.compile(optimizer = 'sgd',
              loss = cat_cross_from_logits,
              metrics = ['categorical_accuracy'])
```

In [19]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 16)	160
conv2d_1 (Conv2D)	(None, 28, 28, 32)	4640
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 5000)	125445000
dense_1 (Dense)	(None, 10)	50010
=====		
Total params:	125,499,810	
Trainable params:	125,499,810	
Non-trainable params:	0	

In [0]:

```
VAL_SPLIT_RATE = 0.1
```

In [0]:

```
EPOCHS_N = 10
```

In [22]:

```
model.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)

Train on 374175 samples, validate on 41576 samples
Epoch 1/10
374175/374175 [=====] - 114s 304us/sample - loss:
0.4815 - categorical_accuracy: 0.8573 - val_loss: 2.9091 - val_categorical_a
ccuracy: 0.3111
Epoch 2/10
374175/374175 [=====] - 114s 304us/sample - loss:
0.3455 - categorical_accuracy: 0.8956 - val_loss: 2.7738 - val_categorical_a
ccuracy: 0.2886
Epoch 3/10
374175/374175 [=====] - 113s 303us/sample - loss:
0.2926 - categorical_accuracy: 0.9110 - val_loss: 1.9089 - val_categorical_a
ccuracy: 0.5299
Epoch 4/10
374175/374175 [=====] - 113s 302us/sample - loss:
0.2517 - categorical_accuracy: 0.9231 - val_loss: 2.2537 - val_categorical_a
ccuracy: 0.4842
Epoch 5/10
374175/374175 [=====] - 114s 303us/sample - loss:
0.2144 - categorical_accuracy: 0.9342 - val_loss: 2.7603 - val_categorical_a
ccuracy: 0.3479
Epoch 6/10
374175/374175 [=====] - 113s 301us/sample - loss:
0.1803 - categorical_accuracy: 0.9447 - val_loss: 2.1086 - val_categorical_a
ccuracy: 0.5454
Epoch 7/10
374175/374175 [=====] - 114s 305us/sample - loss:
0.1466 - categorical_accuracy: 0.9551 - val_loss: 2.4950 - val_categorical_a
ccuracy: 0.4780
Epoch 8/10
374175/374175 [=====] - 113s 302us/sample - loss:
0.1170 - categorical_accuracy: 0.9647 - val_loss: 2.0592 - val_categorical_a
ccuracy: 0.6014
Epoch 9/10
374175/374175 [=====] - 113s 301us/sample - loss:
0.0925 - categorical_accuracy: 0.9728 - val_loss: 2.6105 - val_categorical_a
ccuracy: 0.5198
Epoch 10/10
374175/374175 [=====] - 113s 302us/sample - loss:
0.0738 - categorical_accuracy: 0.9786 - val_loss: 2.7623 - val_categorical_a
ccuracy: 0.5034
```

Out[22]:

```
<tensorflow.python.keras.callbacks.History at 0x7f6ba8e81cf8>
```

In [23]:

```
results = model.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results)
```

```
Test loss, test accuracy: [0.5621323866975915, 0.88420826]
```

Лучшая точность построенной модели на тестовой выборке составила 88%.

## Задание 2

Замените один из сверточных слоев на слой, реализующий операцию пулинга (*Pooling*) с функцией максимума или среднего. Как это повлияло на точность классификатора?

In [0]:

```
from tensorflow.keras.layers import MaxPooling2D

model_2 = tf.keras.Sequential()

model_2.add(Conv2D(16, 3, padding = 'same', activation = 'relu', input_shape = (IMAGE_DIM_0,
model_2.add(MaxPooling2D())
model_2.add(Flatten())
model_2.add(Dense(DENSE_LAYER_WIDTH, activation = 'relu'))
model_2.add(Dense(CLASSES_N))
```

In [0]:

```
model_2.compile(optimizer = 'sgd',
                 loss = cat_crossentropy,
                 metrics = ['categorical_accuracy'])
```

In [26]:

```
model_2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_2 (Conv2D)	(None, 28, 28, 16)	160
=====		
max_pooling2d (MaxPooling2D)	(None, 14, 14, 16)	0
=====		
flatten_1 (Flatten)	(None, 3136)	0
=====		
dense_2 (Dense)	(None, 5000)	15685000
=====		
dense_3 (Dense)	(None, 10)	50010
=====		
Total params: 15,735,170		
Trainable params: 15,735,170		
Non-trainable params: 0		

In [27]:

```
model_2.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)

Train on 374175 samples, validate on 41576 samples
Epoch 1/10
374175/374175 [=====] - 36s 97us/sample - loss: 0.5
802 - categorical_accuracy: 0.8338 - val_loss: 3.2642 - val_categorical_accuracy: 0.0382
Epoch 2/10
374175/374175 [=====] - 36s 96us/sample - loss: 0.4
135 - categorical_accuracy: 0.8762 - val_loss: 3.0404 - val_categorical_accuracy: 0.1671
Epoch 3/10
374175/374175 [=====] - 36s 97us/sample - loss: 0.3
677 - categorical_accuracy: 0.8891 - val_loss: 2.6523 - val_categorical_accuracy: 0.3309
Epoch 4/10
374175/374175 [=====] - 36s 95us/sample - loss: 0.3
384 - categorical_accuracy: 0.8975 - val_loss: 2.3369 - val_categorical_accuracy: 0.4434
Epoch 5/10
374175/374175 [=====] - 36s 96us/sample - loss: 0.3
159 - categorical_accuracy: 0.9037 - val_loss: 2.4419 - val_categorical_accuracy: 0.4272
Epoch 6/10
374175/374175 [=====] - 36s 95us/sample - loss: 0.2
965 - categorical_accuracy: 0.9099 - val_loss: 2.2025 - val_categorical_accuracy: 0.4449
Epoch 7/10
374175/374175 [=====] - 36s 96us/sample - loss: 0.2
797 - categorical_accuracy: 0.9147 - val_loss: 2.2130 - val_categorical_accuracy: 0.4772
Epoch 8/10
374175/374175 [=====] - 37s 98us/sample - loss: 0.2
644 - categorical_accuracy: 0.9193 - val_loss: 2.1537 - val_categorical_accuracy: 0.5218
Epoch 9/10
374175/374175 [=====] - 36s 95us/sample - loss: 0.2
509 - categorical_accuracy: 0.9233 - val_loss: 1.8914 - val_categorical_accuracy: 0.5983
Epoch 10/10
374175/374175 [=====] - 36s 96us/sample - loss: 0.2
375 - categorical_accuracy: 0.9274 - val_loss: 2.2922 - val_categorical_accuracy: 0.5274
```

Out[27]:

```
<tensorflow.python.keras.callbacks.History at 0x7f6ba828da58>
```

In [28]:

```
results_2 = model_2.evaluate(x_test, y_test)
print('Test loss, test accuracy:', results_2)
```

```
Test loss, test accuracy: [0.5017564680594646, 0.87370926]
```

Замена свёрточного слоя на операцию пулинга снизила точность на тестовой выборке до 87%.

## Задание 3

Реализуйте классическую архитектуру сверточных сетей LeNet-5 (<http://yann.lecun.com/exdb/lenet/>) (<http://yann.lecun.com/exdb/lenet/>).

In [0]:

```
from tensorflow.keras.layers import AveragePooling2D

model_3 = tf.keras.Sequential()

model_3.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0, IMAGE_DIM_1, 1)))
model_3.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_3.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0, IMAGE_DIM_1, 1)))
model_3.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_3.add(Flatten())
model_3.add(Dense(120, activation = 'tanh'))
model_3.add(Dense(84, activation = 'tanh'))
model_3.add(Dense(CLASSES_N, activation = 'softmax'))
```

In [0]:

```
model_3.compile(optimizer = 'adam',
                 loss = 'categorical_crossentropy',
                 metrics = ['categorical_accuracy'])
```

In [31]:

```
model_3.fit(x = x, y = y, epochs = EPOCHS_N, validation_split = VAL_SPLIT_RATE)

Train on 374175 samples, validate on 41576 samples
Epoch 1/10
374175/374175 [=====] - 30s 79us/sample - loss: 0.4
504 - categorical_accuracy: 0.8640 - val_loss: 2.7422 - val_categorical_accuracy: 0.2516
Epoch 2/10
374175/374175 [=====] - 30s 80us/sample - loss: 0.3
514 - categorical_accuracy: 0.8912 - val_loss: 2.9686 - val_categorical_accuracy: 0.2343
Epoch 3/10
374175/374175 [=====] - 29s 79us/sample - loss: 0.3
239 - categorical_accuracy: 0.8995 - val_loss: 2.3722 - val_categorical_accuracy: 0.4675
Epoch 4/10
374175/374175 [=====] - 30s 79us/sample - loss: 0.3
086 - categorical_accuracy: 0.9037 - val_loss: 2.9093 - val_categorical_accuracy: 0.3062
Epoch 5/10
374175/374175 [=====] - 30s 79us/sample - loss: 0.2
972 - categorical_accuracy: 0.9069 - val_loss: 2.3950 - val_categorical_accuracy: 0.4489
Epoch 6/10
374175/374175 [=====] - 30s 79us/sample - loss: 0.2
891 - categorical_accuracy: 0.9092 - val_loss: 2.7761 - val_categorical_accuracy: 0.2896
Epoch 7/10
374175/374175 [=====] - 30s 80us/sample - loss: 0.2
811 - categorical_accuracy: 0.9117 - val_loss: 2.7471 - val_categorical_accuracy: 0.3655
Epoch 8/10
374175/374175 [=====] - 30s 80us/sample - loss: 0.2
749 - categorical_accuracy: 0.9136 - val_loss: 1.9750 - val_categorical_accuracy: 0.5368
Epoch 9/10
374175/374175 [=====] - 29s 79us/sample - loss: 0.2
707 - categorical_accuracy: 0.9148 - val_loss: 2.1891 - val_categorical_accuracy: 0.5343
Epoch 10/10
374175/374175 [=====] - 30s 80us/sample - loss: 0.2
659 - categorical_accuracy: 0.9163 - val_loss: 2.2132 - val_categorical_accuracy: 0.5457
```

Out[31]:

```
<tensorflow.python.keras.callbacks.History at 0x7f6ba7593048>
```

In [32]:

```
results_3 = model_3.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results_3)
```

```
Test loss, test accuracy: [0.49347359862197515, 0.8702024]
```

Удивительно, но LeNet-5 показала результат хуже, чем первая модель — 87% на тестовой выборке.

Объяснить это можно тем, что первая модель содержит свёрточный слой с большей выходной размерностью.

## Задание 4

Сравните максимальные точности моделей, построенных в лабораторных работах 1-3. Как можно объяснить полученные различия?

Результаты на валидационной выборке:

- логистическая регрессия — 81%;
- модель с только полно связанными слоями — 10%;
  - с регуляризацией и сбросом нейронов — 72%;
  - с адаптивным шагом — 72%;
- модель с двумя свёрточными слоями и одним полно связанным — 88%;
- модель с одним свёрточным слоем, операцией пулинга и одним полно связанным — 87%;
- *LeNet-5* — два свёрточных слоя две операции пулинга два полно связанных слоя — 87%.

Объяснение превосходства свёрточных сетей над полно связанными — такая архитектура просто предназначена для работы с изображениями.

# Лабораторная работа №4

## Реализация приложения по распознаванию номеров домов

Набор изображений из *Google Street View* с изображениями номеров домов, содержащий 10 классов, соответствующих цифрам от 0 до 9.

- 73257 изображений цифр в обучающей выборке;
- 26032 изображения цифр в тестовой выборке;
- 531131 изображения, которые можно использовать как дополнение к обучающей выборке;
- В двух форматах:
  - Оригинальные изображения с выделенными цифрами;
  - Изображения размером 32×32, содержащие одну цифру;
- Данные первого формата можно скачать по ссылкам:
  - <http://ufldl.stanford.edu/housenumbers/train.tar.gz> (<http://ufldl.stanford.edu/housenumbers/train.tar.gz>) (обучающая выборка);
  - <http://ufldl.stanford.edu/housenumbers/test.tar.gz> (<http://ufldl.stanford.edu/housenumbers/test.tar.gz>) (тестовая выборка);
  - <http://ufldl.stanford.edu/housenumbers/extrtar.gz> (<http://ufldl.stanford.edu/housenumbers/extrtar.gz>) (дополнительные данные);
- Данные второго формата можно скачать по ссылкам:
  - [http://ufldl.stanford.edu/housenumbers/train\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/train_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/train\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/train_32x32.mat)) (обучающая выборка);
  - [http://ufldl.stanford.edu/housenumbers/test\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/test_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/test\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/test_32x32.mat)) (тестовая выборка);
  - [http://ufldl.stanford.edu/housenumbers/extr\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/extr_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/extr\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/extr_32x32.mat)) (дополнительные данные);
- Описание данных на английском языке доступно по ссылке:
  - <http://ufldl.stanford.edu/housenumbers/> (<http://ufldl.stanford.edu/housenumbers/>)

### Задание 1

Реализуйте глубокую нейронную сеть (полносвязную или сверточную) и обучите ее на синтетических данных (например, наборы *MNIST* (<http://yann.lecun.com/exdb/mnist/>) (<http://yann.lecun.com/exdb/mnist/>)) или *notMNIST*).

Ознакомьтесь с имеющимися работами по данной тематике: англоязычная статья (<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf> (<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf>)), видео на *YouTube* ([https://www.youtube.com/watch?v=vGPI\\_JvLoN0](https://www.youtube.com/watch?v=vGPI_JvLoN0)) ([https://www.youtube.com/watch?v=vGPI\\_JvLoN0](https://www.youtube.com/watch?v=vGPI_JvLoN0)).

Используем архитектуру *LeNet-5* и обучим сеть сначала на данных из набора *MNIST*.

In [1]:

```
! pip install tensorflow-gpu --pre --quiet
```

```
! pip show tensorflow-gpu
```

```
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: grpcio, tensorboard, termcolor, six, gast, absl-py, tensorflow-estimator, opt-einsum, wrapt, wheel, numpy, google-pasta, protobuf, scipy, astunparse, keras-preprocessing, h5py
Required-by:
```

In [0]:

```
import tensorflow as tf
from tensorflow import keras
```

In [0]:

```
import numpy as np
```

In [0]:

```
from tensorflow.keras.datasets import mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

In [0]:

```
x_train, x_test = tf.keras.utils.normalize(x_train, axis = 1), tf.keras.utils.normalize(x_t
```

In [0]:

```
x_train, x_test = x_train[..., np.newaxis], x_test[..., np.newaxis]
```

In [7]:

```
from tensorflow.keras.utils import to_categorical
y_train, y_test = to_categorical(y_train), to_categorical(y_test)
y_train.shape
```

Out[7]:

```
(60000, 10)
```

In [0]:

```
IMAGE_DIM_0, IMAGE_DIM_1 = x_train.shape[1], x_train.shape[2]
```

In [0]:

```
CLASSES_N = y_train.shape[1]
```

In [10]:

```
x_train.shape, x_test.shape
```

Out[10]:

```
((60000, 28, 28, 1), (10000, 28, 28, 1))
```

In [0]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import AveragePooling2D, Conv2D, Dense, Flatten

model = tf.keras.Sequential()

model.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
               input_shape = (IMAGE_DIM_0, IMAGE_DIM_1, 1)))
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
               input_shape = (16, 14, 14, 1)))
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Flatten())
model.add(Dense(120, activation = 'tanh'))
model.add(Dense(84, activation = 'tanh'))
model.add(Dense(CLASSES_N, activation = 'softmax'))
```

In [0]:

```
# 'sparse_categorical_crossentropy' gave NAN loss

model.compile(optimizer = 'adam',
              loss = 'categorical_crossentropy',
              metrics = ['categorical_accuracy'])
```

In [13]:

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 6)	156
=====		
average_pooling2d (AveragePo	(None, 14, 14, 6)	0
=====		
conv2d_1 (Conv2D)	(None, 10, 10, 16)	2416
=====		
average_pooling2d_1 (Average	(None, 5, 5, 16)	0
=====		
flatten (Flatten)	(None, 400)	0
=====		
dense (Dense)	(None, 120)	48120
=====		
dense_1 (Dense)	(None, 84)	10164
=====		
dense_2 (Dense)	(None, 10)	850
=====		
Total params:	61,706	
Trainable params:	61,706	
Non-trainable params:	0	

In [0]:

```
EPOCHS_N = 20
```

In [15]:

```
model.fit(x = x_train, y = y_train, validation_split = 0.15, epochs = EPOCHS_N)

Epoch 1/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.2895 - categorical_accuracy: 0.9129 - val_loss: 0.1330 - val_categorical_accuracy: 0.9591
Epoch 2/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.1186 - categorical_accuracy: 0.9639 - val_loss: 0.0995 - val_categorical_accuracy: 0.9694
Epoch 3/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0801 - categorical_accuracy: 0.9748 - val_loss: 0.0960 - val_categorical_accuracy: 0.9697
Epoch 4/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0631 - categorical_accuracy: 0.9799 - val_loss: 0.0737 - val_categorical_accuracy: 0.9786
Epoch 5/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0484 - categorical_accuracy: 0.9848 - val_loss: 0.0679 - val_categorical_accuracy: 0.9802
Epoch 6/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0394 - categorical_accuracy: 0.9874 - val_loss: 0.0647 - val_categorical_accuracy: 0.9806
Epoch 7/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0328 - categorical_accuracy: 0.9894 - val_loss: 0.0688 - val_categorical_accuracy: 0.9811
Epoch 8/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0268 - categorical_accuracy: 0.9919 - val_loss: 0.0698 - val_categorical_accuracy: 0.9806
Epoch 9/20
1594/1594 [=====] - 4s 3ms/step - loss: 0.0242 - categorical_accuracy: 0.9921 - val_loss: 0.0617 - val_categorical_accuracy: 0.9822
Epoch 10/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0197 - categorical_accuracy: 0.9935 - val_loss: 0.0687 - val_categorical_accuracy: 0.9822
Epoch 11/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0168 - categorical_accuracy: 0.9944 - val_loss: 0.0719 - val_categorical_accuracy: 0.9811
Epoch 12/20
1594/1594 [=====] - 4s 3ms/step - loss: 0.0155 - categorical_accuracy: 0.9948 - val_loss: 0.0680 - val_categorical_accuracy: 0.9808
Epoch 13/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0137 - categorical_accuracy: 0.9955 - val_loss: 0.0679 - val_categorical_accuracy: 0.9822
Epoch 14/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0121 - categorical_accuracy: 0.9965 - val_loss: 0.0768 - val_categorical_accuracy: 0.9806
Epoch 15/20
```

```
1594/1594 [=====] - 5s 3ms/step - loss: 0.0115 - categorical_accuracy: 0.9961 - val_loss: 0.0747 - val_categorical_accuracy: 0.9824
Epoch 16/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0122 - categorical_accuracy: 0.9958 - val_loss: 0.0687 - val_categorical_accuracy: 0.9840
Epoch 17/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0095 - categorical_accuracy: 0.9966 - val_loss: 0.0633 - val_categorical_accuracy: 0.9848
Epoch 18/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0100 - categorical_accuracy: 0.9965 - val_loss: 0.0701 - val_categorical_accuracy: 0.9836
Epoch 19/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0055 - categorical_accuracy: 0.9984 - val_loss: 0.0706 - val_categorical_accuracy: 0.9842
Epoch 20/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0093 - categorical_accuracy: 0.9967 - val_loss: 0.0714 - val_categorical_accuracy: 0.9838
```

Out[15]:

```
<tensorflow.python.keras.callbacks.History at 0x7f60025300f0>
```

In [16]:

```
results = model.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.0828 - categorical_accuracy: 0.9808
Test loss, test accuracy: [0.08280119299888611, 0.9807999730110168]
```

Удалось достичь отличного результата — точность распознавания на тестовой выборке составила 98,0%.

## Задание 2

После уточнения модели на синтетических данных попробуйте обучить ее на реальных данных (набор *Google Street View*). Что изменилось в модели?

In [0]:

```
DS_URL_FOLDER = 'http://ufldl.stanford.edu/housenumbers/'

FIRST_DS_EXT = '.tar.gz'
SECOND_DS_EXT = '_32x32.mat'

TRAIN_DS_NAME = 'train'
TEST_DS_NAME = 'test'
EXTRA_DS_NAME = 'extra'
```

In [0]:

```
from urllib.request import urlretrieve
import tarfile
import os

def load_file(_url_folder, _name, _ext, _key, _local_ext = ''):

    file_url_ = _url_folder + _name + _ext

    local_file_name_ = _name + '_' + _key + _local_ext

    urlretrieve(file_url_, local_file_name_)

    return local_file_name_

def tar_gz_to_dir(_url_folder, _name, _ext, _key):

    local_file_name_ = load_file(_url_folder, _name, _ext, _key, _ext)

    dir_name_ = _name + '_' + _key

    with tarfile.open(local_file_name_, 'r:gz') as tar_:
        tar_.extractall(dir_name_)

    os.remove(local_file_name_)

    return dir_name_
```

In [0]:

```
first_ds_train_dir = tar_gz_to_dir(DS_URL_FOLDER, TRAIN_DS_NAME, FIRST_DS_EXT, 'first')
first_ds_test_dir = tar_gz_to_dir(DS_URL_FOLDER, TEST_DS_NAME, FIRST_DS_EXT, 'first')
first_ds_extra_dir = tar_gz_to_dir(DS_URL_FOLDER, EXTRA_DS_NAME, FIRST_DS_EXT, 'first')
```

In [0]:

```
second_ds_train_file = load_file(DS_URL_FOLDER, TRAIN_DS_NAME, SECOND_DS_EXT, 'second')
second_ds_test_file = load_file(DS_URL_FOLDER, TEST_DS_NAME, SECOND_DS_EXT, 'second')
second_ds_extra_file = load_file(DS_URL_FOLDER, EXTRA_DS_NAME, SECOND_DS_EXT, 'second')
```

In [0]:

```
from scipy import io

second_ds_train = io.loadmat(second_ds_train_file)
second_ds_test = io.loadmat(second_ds_test_file)
second_ds_extra = io.loadmat(second_ds_extra_file)
```

In [22]:

```
X_second_ds_train = np.moveaxis(second_ds_train['X'], -1, 0)
X_second_ds_test = np.moveaxis(second_ds_test['X'], -1, 0)
X_second_ds_extra = np.moveaxis(second_ds_extra['X'], -1, 0)

y_second_ds_train = second_ds_train['y']
y_second_ds_test = second_ds_test['y']
y_second_ds_extra = second_ds_extra['y']

print(X_second_ds_train.shape, y_second_ds_train.shape)
print(X_second_ds_test.shape, y_second_ds_test.shape)
print(X_second_ds_extra.shape, y_second_ds_extra.shape)
```

```
(73257, 32, 32, 3) (73257, 1)
(26032, 32, 32, 3) (26032, 1)
(531131, 32, 32, 3) (531131, 1)
```

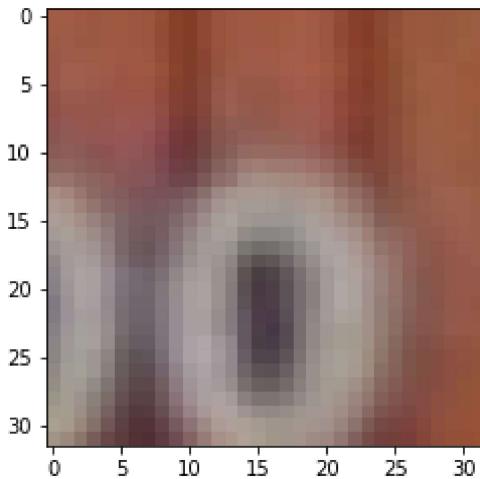
In [23]:

```
import matplotlib.pyplot as plt

plt.imshow(X_second_ds_train[100])
plt.imshow(X_second_ds_test[100])
plt.imshow(X_second_ds_extra[100])
```

Out[23]:

```
<matplotlib.image.AxesImage at 0x7f6002399320>
```



In [0]:

```
IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2 = X_second_ds_train.shape[-3], X_second_ds_train
```

In [0]:

```
y_second_ds_train_cat = to_categorical(y_second_ds_train)
y_second_ds_test_cat = to_categorical(y_second_ds_test)
```

In [0]:

```
CLASSES_N_2 = y_second_ds_train_cat.shape[1]
```

In [0]:

```
model_2 = tf.keras.Sequential()

model_2.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2)))
model_2.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_2.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2)))
model_2.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_2.add(Flatten())
model_2.add(Dense(120, activation = 'tanh'))
model_2.add(Dense(84, activation = 'tanh'))
model_2.add(Dense(CLASSES_N_2, activation = 'softmax'))
```

In [0]:

```
model_2.compile(optimizer = 'adam',
                 loss = 'categorical_crossentropy',
                 metrics = ['categorical_accuracy'])
```

In [29]:

```
model_2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_2 (Conv2D)	(None, 32, 32, 6)	456
<hr/>		
average_pooling2d_2 (Average)	(None, 16, 16, 6)	0
<hr/>		
conv2d_3 (Conv2D)	(None, 12, 12, 16)	2416
<hr/>		
average_pooling2d_3 (Average)	(None, 6, 6, 16)	0
<hr/>		
flatten_1 (Flatten)	(None, 576)	0
<hr/>		
dense_3 (Dense)	(None, 120)	69240
<hr/>		
dense_4 (Dense)	(None, 84)	10164
<hr/>		
dense_5 (Dense)	(None, 11)	935
<hr/>		
Total params: 83,211		
Trainable params: 83,211		
Non-trainable params: 0		

In [30]:

```
model_2.fit(x = X_second_ds_train, y = y_second_ds_train_cat, validation_split = 0.15, epoch
```

Epoch 1/20  
1946/1946 [=====] - 6s 3ms/step - loss: 1.1878 - categorical\_accuracy: 0.6141 - val\_loss: 0.8286 - val\_categorical\_accuracy: 0.7351  
Epoch 2/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.6917 - categorical\_accuracy: 0.7843 - val\_loss: 0.6128 - val\_categorical\_accuracy: 0.8132  
Epoch 3/20  
1946/1946 [=====] - 6s 3ms/step - loss: 0.5979 - categorical\_accuracy: 0.8130 - val\_loss: 0.5611 - val\_categorical\_accuracy: 0.8280  
Epoch 4/20  
1946/1946 [=====] - 6s 3ms/step - loss: 0.5475 - categorical\_accuracy: 0.8287 - val\_loss: 0.5569 - val\_categorical\_accuracy: 0.8276  
Epoch 5/20  
1946/1946 [=====] - 6s 3ms/step - loss: 0.5193 - categorical\_accuracy: 0.8383 - val\_loss: 0.5476 - val\_categorical\_accuracy: 0.8333  
Epoch 6/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.4886 - categorical\_accuracy: 0.8474 - val\_loss: 0.5300 - val\_categorical\_accuracy: 0.8375  
Epoch 7/20  
1946/1946 [=====] - 6s 3ms/step - loss: 0.4755 - categorical\_accuracy: 0.8511 - val\_loss: 0.5751 - val\_categorical\_accuracy: 0.8284  
Epoch 8/20  
1946/1946 [=====] - 6s 3ms/step - loss: 0.4482 - categorical\_accuracy: 0.8613 - val\_loss: 0.5043 - val\_categorical\_accuracy: 0.8469  
Epoch 9/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.4372 - categorical\_accuracy: 0.8642 - val\_loss: 0.5296 - val\_categorical\_accuracy: 0.8393  
Epoch 10/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.4162 - categorical\_accuracy: 0.8697 - val\_loss: 0.5304 - val\_categorical\_accuracy: 0.8381  
Epoch 11/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.4058 - categorical\_accuracy: 0.8733 - val\_loss: 0.5445 - val\_categorical\_accuracy: 0.8385  
Epoch 12/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.4027 - categorical\_accuracy: 0.8753 - val\_loss: 0.5371 - val\_categorical\_accuracy: 0.8347  
Epoch 13/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.3966 - categorical\_accuracy: 0.8747 - val\_loss: 0.5472 - val\_categorical\_accuracy: 0.8381  
Epoch 14/20  
1946/1946 [=====] - 5s 3ms/step - loss: 0.3922 - categorical\_accuracy: 0.8759 - val\_loss: 0.5105 - val\_categorical\_accuracy: 0.8482  
Epoch 15/20

```
1946/1946 [=====] - 5s 3ms/step - loss: 0.3835 - categorical_accuracy: 0.8784 - val_loss: 0.5198 - val_categorical_accuracy: 0.8477
Epoch 16/20
1946/1946 [=====] - 5s 3ms/step - loss: 0.3669 - categorical_accuracy: 0.8850 - val_loss: 0.4903 - val_categorical_accuracy: 0.8523
Epoch 17/20
1946/1946 [=====] - 5s 3ms/step - loss: 0.3616 - categorical_accuracy: 0.8858 - val_loss: 0.5543 - val_categorical_accuracy: 0.8359
Epoch 18/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.3559 - categorical_accuracy: 0.8873 - val_loss: 0.5141 - val_categorical_accuracy: 0.8467
Epoch 19/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.3472 - categorical_accuracy: 0.8914 - val_loss: 0.5160 - val_categorical_accuracy: 0.8511
Epoch 20/20
1946/1946 [=====] - 5s 3ms/step - loss: 0.3439 - categorical_accuracy: 0.8917 - val_loss: 0.6758 - val_categorical_accuracy: 0.7925
```

Out[30]:

```
<tensorflow.python.keras.callbacks.History at 0x7f60022ca4e0>
```

In [31]:

```
results = model_2.evaluate(X_second_ds_test, y_second_ds_test_cat)
print('Test loss, test accuracy:', results)
```

```
814/814 [=====] - 1s 2ms/step - loss: 0.7207 - categorical_accuracy: 0.7746
Test loss, test accuracy: [0.7207155227661133, 0.7745851278305054]
```

Прежде всего, в модели изменилось то, что добавился ещё один класс — *не распознано*.

Эти данные более сложны для распознавания, что повлияло на результат — точность распознавания на тестовой выборке составила 77,4%.

## Задание 3

Сделайте множество снимков изображений номеров домов с помощью смартфона на ОС *Android*. Также можно использовать библиотеки *OpenCV*, *Simple CV* или *Pygame* для обработки изображений с общедоступных камер видеонаблюдения (например, <https://www.earthcam.com/>) (<https://www.earthcam.com/>)).

В качестве примера использования библиотеки *TensorFlow* на смартфоне можете воспользоваться демонстрационным приложением от *Google* (<https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/android>) (<https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/android>)).

## Задание 4

Реализуйте приложение для ОС *Android*, которое может распознавать цифры в номерах домов, используя разработанный ранее классификатор. Какова доля правильных классификаций?

# Лабораторная работа №5

## Применение сверточных нейронных сетей (бинарная классификация)

Набор данных *DogsVsCats*, который состоит из изображений различной размерности, содержащих фотографии собак и кошек.

Обучающая выборка включает в себя 25 тыс. изображений (12,5 тыс. кошек: *cat.0.jpg*, ..., *cat.12499.jpg* и 12,5 тыс. собак: *dog.0.jpg*, ..., *dog.12499.jpg*), а контрольная выборка содержит 12,5 тыс. неразмеченных изображений.

Скачать данные, а также проверить качество классификатора на тестовой выборке можно на сайте Kaggle: <https://www.kaggle.com/c/dogs-vs-cats/data> (<https://www.kaggle.com/c/dogs-vs-cats/data>)

### Задание 1

Загрузите данные. Разделите исходный набор данных на обучающую, валидационную и контрольную выборки.

In [1]:

```
from google.colab import drive  
drive.mount('/content/drive', force_remount = True)
```

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2/dogs-vs-cats'  
  
import sys  
  
sys.path.append(BASE_DIR)  
  
import os
```

In [0]:

```
TRAIN_ARCHIVE_NAME = 'train.zip'  
TEST_ARCHIVE_NAME = 'test1.zip'  
  
LOCAL_DIR_NAME = 'dogs-vs-cats'
```

In [0]:

```
from zipfile import ZipFile

with ZipFile(os.path.join(BASE_DIR, TRAIN_ARCHIVE_NAME), 'r') as zip_:
    zip_.extractall(path = os.path.join(LOCAL_DIR_NAME, 'train'))

with ZipFile(os.path.join(BASE_DIR, TEST_ARCHIVE_NAME), 'r') as zip_:
    zip_.extractall(path = os.path.join(LOCAL_DIR_NAME, 'test-1'))
```

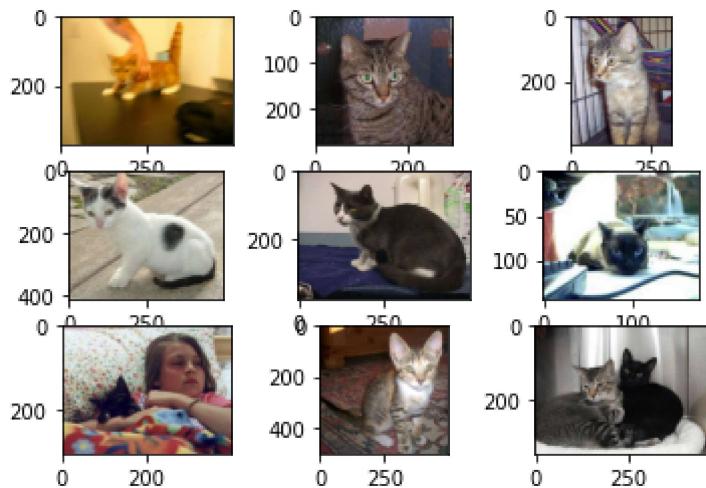
In [5]:

```
from matplotlib import pyplot
from matplotlib.image import imread

dir_ = 'dogs-vs-cats/train/train'

for i in range(9):
    pyplot.subplot(330 + 1 + i)
    image_ = imread('{}/cat.{}.jpg'.format(dir_, i))
    pyplot.imshow(image_)

pyplot.show()
```



In [0]:

```
NEW_IMAGE_WIDTH = 100
```

In [7]:

```
from os import listdir
from os.path import join
from numpy import asarray
from numpy import save
from keras.preprocessing.image import load_img
from keras.preprocessing.image import img_to_array

def dir_to_dataset(_dir_path):

    photos_, labels_ = [], []
    for file_ in listdir(_dir_path):
        if file_.startswith('cat'):
            label_ = 1.0
        else:
            label_ = 0.0
        photo_ = load_img(join(_dir_path, file_), target_size = (NEW_IMAGE_WIDTH, NEW_IMAGE_HEIGHT))
        photo_ = img_to_array(photo_)
        photos_.append(photo_)
        labels_.append(label_)

    photos_norm_ = tf.keras.utils.normalize(photos_, axis = 1)
    return asarray(photos_norm_), asarray(labels_)
```

Using TensorFlow backend.

In [8]:

```
! pip install tensorflow-gpu --pre --quiet
! pip show tensorflow-gpu
```

Name: tensorflow-gpu  
Version: 2.2.0rc3  
Summary: TensorFlow is an open source machine learning framework for everyone.  
Home-page: <https://www.tensorflow.org/> (<https://www.tensorflow.org/>)  
Author: Google Inc.  
Author-email: packages@tensorflow.org  
License: Apache 2.0  
Location: /usr/local/lib/python3.6/dist-packages  
Requires: grpcio, h5py, gast, numpy, protobuf, google-pasta, absl-py, astunparse, opt-einsum, six, wheel, scipy, tensorflow-estimator, keras-preprocessing, tensorboard, termcolor, wrapt  
Required-by:

In [0]:

```
import tensorflow as tf
```

```
In [0]:
```

```
import numpy as np
```

```
In [0]:
```

```
X_all, y_all = dir_to_dataset('dogs-vs-cats/train/train')
```

```
In [0]:
```

```
TEST_LEN_HALF = 1000
```

```
In [13]:
```

```
test_interval = np.r_[0:TEST_LEN_HALF, -TEST_LEN_HALF:-0]  
X, y = X_all[TEST_LEN_HALF:-TEST_LEN_HALF], y_all[TEST_LEN_HALF:-TEST_LEN_HALF]  
X_test, y_test = X_all[test_interval], y_all[test_interval]  
  
print(X.shape, y.shape)  
print(X_test.shape, y_test.shape)
```

```
(23000, 100, 100, 3) (23000,)  
(2000, 100, 100, 3) (2000,)
```

Выделение валидационной выборки произойдёт автоматически по параметру `validation_split` метода `model.fit()`.

## Задание 2

Реализуйте глубокую нейронную сеть с как минимум тремя сверточными слоями. Какое качество классификации получено?

```
In [0]:
```

```
from tensorflow import keras
```

In [15]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense

model = tf.keras.Sequential()

model.add(Conv2D(16, 3, padding = 'same', activation = 'relu', input_shape = (NEW_IMAGE_WID
model.add(MaxPooling2D())
model.add(Conv2D(32, 3, padding = 'same', activation = 'relu'))
model.add(MaxPooling2D())
model.add(Conv2D(64, 3, padding = 'same', activation = 'relu'))
model.add(MaxPooling2D())
model.add(Flatten())
model.add(Dense(512, activation = 'relu'))
model.add(Dense(1, activation = 'sigmoid'))

model.compile(optimizer = 'sgd',
              loss = 'binary_crossentropy',
              metrics = ['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
<hr/>		
conv2d (Conv2D)	(None, 100, 100, 16)	448
<hr/>		
max_pooling2d (MaxPooling2D)	(None, 50, 50, 16)	0
<hr/>		
conv2d_1 (Conv2D)	(None, 50, 50, 32)	4640
<hr/>		
max_pooling2d_1 (MaxPooling2D)	(None, 25, 25, 32)	0
<hr/>		
conv2d_2 (Conv2D)	(None, 25, 25, 64)	18496
<hr/>		
max_pooling2d_2 (MaxPooling2D)	(None, 12, 12, 64)	0
<hr/>		
flatten (Flatten)	(None, 9216)	0
<hr/>		
dense (Dense)	(None, 512)	4719104
<hr/>		
dense_1 (Dense)	(None, 1)	513
<hr/>		
Total params: 4,743,201		
Trainable params: 4,743,201		
Non-trainable params: 0		

In [16]:

```
model.fit(x = X, y = y, epochs = 20, validation_split = 0.15)
```

```
Epoch 1/20
611/611 [=====] - 5s 8ms/step - loss: 0.6907 - accuracy: 0.5365 - val_loss: 0.6885 - val_accuracy: 0.5670
Epoch 2/20
611/611 [=====] - 4s 7ms/step - loss: 0.6858 - accuracy: 0.5711 - val_loss: 0.6845 - val_accuracy: 0.5481
Epoch 3/20
611/611 [=====] - 4s 7ms/step - loss: 0.6794 - accuracy: 0.5812 - val_loss: 0.6754 - val_accuracy: 0.5957
Epoch 4/20
611/611 [=====] - 4s 7ms/step - loss: 0.6724 - accuracy: 0.5894 - val_loss: 0.6685 - val_accuracy: 0.5875
Epoch 5/20
611/611 [=====] - 4s 7ms/step - loss: 0.6626 - accuracy: 0.6043 - val_loss: 0.6565 - val_accuracy: 0.6104
Epoch 6/20
611/611 [=====] - 4s 7ms/step - loss: 0.6529 - accuracy: 0.6184 - val_loss: 0.6447 - val_accuracy: 0.6301
Epoch 7/20
611/611 [=====] - 4s 7ms/step - loss: 0.6425 - accuracy: 0.6298 - val_loss: 0.6570 - val_accuracy: 0.6038
Epoch 8/20
611/611 [=====] - 4s 7ms/step - loss: 0.6313 - accuracy: 0.6492 - val_loss: 0.6229 - val_accuracy: 0.6600
Epoch 9/20
611/611 [=====] - 4s 7ms/step - loss: 0.6184 - accuracy: 0.6617 - val_loss: 0.6195 - val_accuracy: 0.6501
Epoch 10/20
611/611 [=====] - 4s 7ms/step - loss: 0.6086 - accuracy: 0.6730 - val_loss: 0.6173 - val_accuracy: 0.6562
Epoch 11/20
611/611 [=====] - 4s 7ms/step - loss: 0.5992 - accuracy: 0.6808 - val_loss: 0.6338 - val_accuracy: 0.6336
Epoch 12/20
611/611 [=====] - 4s 7ms/step - loss: 0.5900 - accuracy: 0.6877 - val_loss: 0.6000 - val_accuracy: 0.6771
Epoch 13/20
611/611 [=====] - 4s 7ms/step - loss: 0.5791 - accuracy: 0.6984 - val_loss: 0.5896 - val_accuracy: 0.6829
Epoch 14/20
611/611 [=====] - 4s 7ms/step - loss: 0.5655 - accuracy: 0.7121 - val_loss: 0.5739 - val_accuracy: 0.6980
Epoch 15/20
611/611 [=====] - 4s 7ms/step - loss: 0.5529 - accuracy: 0.7184 - val_loss: 0.5741 - val_accuracy: 0.7041
Epoch 16/20
611/611 [=====] - 4s 7ms/step - loss: 0.5407 - accuracy: 0.7278 - val_loss: 0.5559 - val_accuracy: 0.7113
Epoch 17/20
611/611 [=====] - 4s 7ms/step - loss: 0.5239 - accuracy: 0.7384 - val_loss: 0.5465 - val_accuracy: 0.7235
Epoch 18/20
611/611 [=====] - 4s 7ms/step - loss: 0.5113 - accuracy: 0.7481 - val_loss: 0.5411 - val_accuracy: 0.7365
Epoch 19/20
611/611 [=====] - 4s 7ms/step - loss: 0.4984 - accuracy: 0.7573 - val_loss: 0.5382 - val_accuracy: 0.7374
```

```
Epoch 20/20
611/611 [=====] - 4s 7ms/step - loss: 0.4847 - accuracy: 0.7671 - val_loss: 0.5281 - val_accuracy: 0.7403
```

Out[16]:

```
<tensorflow.python.keras.callbacks.History at 0x7f31f0e59198>
```

In [17]:

```
results = model.evaluate(X_test, y_test)
print('Test loss, test accuracy:', results)
```

```
63/63 [=====] - 0s 4ms/step - loss: 0.5323 - accuracy: 0.7260
Test loss, test accuracy: [0.5323428511619568, 0.7260000109672546]
```

Результат — 72% на тестовой выборке.

### Задание 3

Примените дополнение данных (*data augmentation*). Как это повлияло на качество классификатора?

### Задание 4

Поэкспериментируйте с готовыми нейронными сетями (например, *AlexNet*, *VGG16*, *Inception* и т.п.), применив передаточное обучение. Как это повлияло на качество классификатора?

Какой максимальный результат удалось получить на сайте *Kaggle*? Почему?

# Лабораторная работа №6

## Применение сверточных нейронных сетей (многоклассовая классификация)

Набор данных для распознавания языка жестов, который состоит из изображений размерности 28x28 в оттенках серого (значение пикселя от 0 до 255).

Каждое из изображений обозначает букву латинского алфавита, обозначенную с помощью жеста (изображения в наборе данных в оттенках серого).

Обучающая выборка включает в себя 27,455 изображений, а контрольная выборка содержит 7172 изображения.

Данные в виде csv-файлов можно скачать на сайте Kaggle: <https://www.kaggle.com/datamunge/sign-language-mnist> (<https://www.kaggle.com/datamunge/sign-language-mnist>)

### Задание 1

Загрузите данные. Разделите исходный набор данных на обучающую и валидационную выборки.

In [1]:

```
from google.colab import drive  
  
drive.mount('/content/drive', force_remount = True)
```

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2'  
  
import sys  
  
sys.path.append(BASE_DIR)  
  
import os
```

In [0]:

```
DATA_ARCHIVE_NAME = 'sign-language-mnist.zip'  
  
LOCAL_DIR_NAME = 'sign-language'
```

In [0]:

```
from zipfile import ZipFile  
  
with ZipFile(os.path.join(BASE_DIR, DATA_ARCHIVE_NAME), 'r') as zip_:  
    zip_.extractall(path = os.path.join(LOCAL_DIR_NAME, 'train'))
```

In [0]:

```
TRAIN_FILE_PATH = 'sign-language/train/sign_mnist_train.csv'  
TEST_FILE_PATH = 'sign-language/train/sign_mnist_test.csv'
```

In [0]:

```
import pandas as pd  
  
train_df = pd.read_csv(TRAIN_FILE_PATH)  
test_df = pd.read_csv(TEST_FILE_PATH)
```

In [7]:

```
train_df.shape, test_df.shape
```

Out[7]:

```
((27455, 785), (7172, 785))
```

In [0]:

```
IMAGE_DIM = 28
```

In [0]:

```
def row_to_label(_row):  
    return _row[0]  
  
def row_to_one_image(_row):  
    return _row[1:].values.reshape((IMAGE_DIM, IMAGE_DIM, 1))
```

In [0]:

```
def to_images_and_labels(_dataframe):  
  
    llll = _dataframe.apply(lambda row: row_to_label(row), axis = 1)  
    mmmm = _dataframe.apply(lambda row: row_to_one_image(row), axis = 1)  
  
    data_dict_ = { 'label': llll, 'image': mmmm }  
  
    reshaped_ = pd.DataFrame(data_dict_, columns = ['label', 'image'])  
  
    return reshaped_
```

In [0]:

```
train_df_reshaped = to_images_and_labels(train_df)  
test_df_reshaped = to_images_and_labels(test_df)
```

## Задание 2

Реализуйте глубокую нейронную сеть со сверточными слоями. Какое качество классификации получено? Какая архитектура сети была использована?

Возьмём LeNet-5.

In [12]:

```
! pip install tensorflow-gpu --pre --quiet
```

```
! pip show tensorflow-gpu
```

```
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: keras-preprocessing, termcolor, absl-py, grpcio, wheel, scipy, astunparse, numpy, tensorflow-estimator, h5py, protobuf, opt-einsum, gast, six, tensorflowboard, wrapt, google-pasta
Required-by:
```

In [0]:

```
from tensorflow.keras.utils import to_categorical
import numpy as np

X_train = np.asarray(list(train_df_reshaped['image']))
X_test = np.asarray(list(test_df_reshaped['image']))

y_train = to_categorical(train_df_reshaped['label'].astype('category').cat.codes.astype('int32'))
y_test = to_categorical(test_df_reshaped['label'].astype('category').cat.codes.astype('int32'))
```

In [14]:

```
X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

Out[14]:

```
((27455, 28, 28, 1), (27455, 24), (7172, 28, 28, 1), (7172, 24))
```

In [0]:

```
CLASSES_N = y_train.shape[1]
```

In [0]:

```
import tensorflow as tf
```

In [0]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import AveragePooling2D, Conv2D, Dense, Flatten

model = tf.keras.Sequential()

model.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
               input_shape = (IMAGE_DIM, IMAGE_DIM, 1)))
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
               input_shape = (IMAGE_DIM // 2, IMAGE_DIM // 2, 1)))
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Flatten())
model.add(Dense(120, activation = 'tanh'))
model.add(Dense(84, activation = 'tanh'))
model.add(Dense(CLASSES_N, activation = 'softmax'))
```

In [0]:

```
model.compile(optimizer = 'adam',
              loss = 'categorical_crossentropy',
              metrics = ['categorical_accuracy'])
```

In [19]:

```
model.fit(x = X_train, y = y_train, epochs = 20, validation_split = 0.15)
```

Epoch 1/20

730/730 [=====] - 3s 3ms/step - loss: 0.7144 - categorical\_accuracy: 0.8216 - val\_loss: 0.1296 - val\_categorical\_accuracy: 0.9857

Epoch 2/20

730/730 [=====] - 2s 3ms/step - loss: 0.0620 - categorical\_accuracy: 0.9943 - val\_loss: 0.0209 - val\_categorical\_accuracy: 0.9995

Epoch 3/20

730/730 [=====] - 2s 3ms/step - loss: 0.0136 - categorical\_accuracy: 1.0000 - val\_loss: 0.0083 - val\_categorical\_accuracy: 0.9995

Epoch 4/20

730/730 [=====] - 2s 3ms/step - loss: 0.1373 - categorical\_accuracy: 0.9634 - val\_loss: 0.0284 - val\_categorical\_accuracy: 0.9976

Epoch 5/20

730/730 [=====] - 2s 3ms/step - loss: 0.0107 - categorical\_accuracy: 0.9998 - val\_loss: 0.0054 - val\_categorical\_accuracy: 1.0000

Epoch 6/20

730/730 [=====] - 2s 3ms/step - loss: 0.0036 - categorical\_accuracy: 1.0000 - val\_loss: 0.0028 - val\_categorical\_accuracy: 1.0000

Epoch 7/20

730/730 [=====] - 2s 3ms/step - loss: 0.0020 - categorical\_accuracy: 1.0000 - val\_loss: 0.0018 - val\_categorical\_accuracy: 1.0000

Epoch 8/20

730/730 [=====] - 2s 3ms/step - loss: 0.0013 - categorical\_accuracy: 1.0000 - val\_loss: 0.0013 - val\_categorical\_accuracy: 1.0000

Epoch 9/20

730/730 [=====] - 2s 3ms/step - loss: 8.6329e-04 - categorical\_accuracy: 1.0000 - val\_loss: 8.2630e-04 - val\_categorical\_accuracy: 1.0000

Epoch 10/20

730/730 [=====] - 2s 3ms/step - loss: 5.9700e-04 - categorical\_accuracy: 1.0000 - val\_loss: 5.6174e-04 - val\_categorical\_accuracy: 1.0000

Epoch 11/20

730/730 [=====] - 2s 3ms/step - loss: 3.8820e-04 - categorical\_accuracy: 1.0000 - val\_loss: 3.7687e-04 - val\_categorical\_accuracy: 1.0000

Epoch 12/20

730/730 [=====] - 2s 3ms/step - loss: 2.6361e-04 - categorical\_accuracy: 1.0000 - val\_loss: 2.7043e-04 - val\_categorical\_accuracy: 1.0000

Epoch 13/20

730/730 [=====] - 2s 3ms/step - loss: 1.8208e-04 - categorical\_accuracy: 1.0000 - val\_loss: 1.9701e-04 - val\_categorical\_accuracy: 1.0000

Epoch 14/20

730/730 [=====] - 2s 3ms/step - loss: 1.2523e-04 - categorical\_accuracy: 1.0000 - val\_loss: 1.3436e-04 - val\_categorical\_accuracy: 1.0000

Epoch 15/20

```
730/730 [=====] - 2s 3ms/step - loss: 8.3785e-05 -  
categorical_accuracy: 1.0000 - val_loss: 9.1194e-05 - val_categorical_accuracy: 1.0000  
Epoch 16/20  
730/730 [=====] - 2s 3ms/step - loss: 5.6255e-05 -  
categorical_accuracy: 1.0000 - val_loss: 8.4611e-05 - val_categorical_accuracy: 1.0000  
Epoch 17/20  
730/730 [=====] - 2s 3ms/step - loss: 3.8880e-05 -  
categorical_accuracy: 1.0000 - val_loss: 4.7102e-05 - val_categorical_accuracy: 1.0000  
Epoch 18/20  
730/730 [=====] - 2s 3ms/step - loss: 0.2848 - categorical_accuracy: 0.9219 - val_loss: 0.1512 - val_categorical_accuracy: 0.9641  
Epoch 19/20  
730/730 [=====] - 2s 3ms/step - loss: 0.0610 - categorical_accuracy: 0.9873 - val_loss: 0.0155 - val_categorical_accuracy: 0.9995  
Epoch 20/20  
730/730 [=====] - 2s 3ms/step - loss: 0.0104 - categorical_accuracy: 0.9992 - val_loss: 0.0050 - val_categorical_accuracy: 0.9995
```

Out[19]:

```
<tensorflow.python.keras.callbacks.History at 0x7f51b04778d0>
```

In [20]:

```
results = model.evaluate(X_test, y_test)  
print('Test loss, test accuracy:', results)
```

```
225/225 [=====] - 0s 2ms/step - loss: 0.3798 - categorical_accuracy: 0.8755  
Test loss, test accuracy: [0.37977278232574463, 0.8754879832267761]
```

За 20 эпох удалось достичь точности в 87% на тестовой выборке.

## Задание 3

Примените дополнение данных (*data augmentation*). Как это повлияло на качество классификатора?

## Задание 4

Поэкспериментируйте с готовыми нейронными сетями (например, *AlexNet*, *VGG16*, *Inception* и т.п.), применив передаточное обучение. Как это повлияло на качество классификатора? Можно ли было обойтись без него?

Какой максимальный результат удалось получить на контрольной выборке?

# Лабораторная работа №7

## Рекуррентные нейронные сети для анализа текста

Набор данных для предсказания оценок для отзывов, собранных с сайта *imdb.com*, который состоит из 50,000 отзывов в виде текстовых файлов.

Отзывы разделены на положительные (25,000) и отрицательные (25,000).

Данные предварительно токенизированы по принципу «мешка слов», индексы слов можно взять из словаря (*imdb.vocab*).

Обучающая выборка включает в себя 12,500 положительных и 12,500 отрицательных отзывов, контрольная выборка также содержит 12,500 положительных и 12,500 отрицательных отзывов.

Данные можно скачать на сайте Kaggle: <https://www.kaggle.com/iarunava/imdb-movie-reviews-dataset> (<https://www.kaggle.com/iarunava/imdb-movie-reviews-dataset>) <https://ai.stanford.edu/~amaas/data/sentiment/> (<https://ai.stanford.edu/~amaas/data/sentiment/>)

### Задание 1

Загрузите данные. Преобразуйте текстовые файлы во внутренние структуры данных, которые используют индексы вместо слов.

Будем брать первые MAX\_LENGTH слов, а если в отзыве слов меньше, чем это число, то применять паддинг.

In [1]:

```
from google.colab import drive  
  
drive.mount('/content/drive', force_remount = True)
```

Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response\\_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly \(https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect\\_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response\\_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly\)](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly (https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aoob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly))

Enter your authorization code:

.....

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2'

import sys

sys.path.append(BASE_DIR)

import os
```

In [0]:

```
DATA_ARCHIVE_NAME = 'imdb-dataset-of-50k-movie-reviews.zip'

LOCAL_DIR_NAME = 'imdb-sentiments'
```

In [0]:

```
from zipfile import ZipFile

with ZipFile(os.path.join(BASE_DIR, DATA_ARCHIVE_NAME), 'r') as zip_:
    zip_.extractall(LOCAL_DIR_NAME)
```

In [0]:

```
DATA_FILE_PATH = 'imdb-sentiments/IMDB Dataset.csv'
```

In [0]:

```
import pandas as pd

all_df = pd.read_csv(DATA_FILE_PATH)
```

In [0]:

```
df_test = all_df.sample(frac = 0.1)

df_train = all_df.drop(df_test.index)
```

In [8]:

```
df_train.shape, df_test.shape
```

Out[8]:

```
((45000, 2), (5000, 2))
```

In [9]:

```
import nltk  
  
nltk.download('punkt')
```

[nltk\_data] Downloading package punkt to /root/nltk\_data...  
[nltk\_data] Unzipping tokenizers/punkt.zip.

Out[9]:

True

In [0]:

```
MAX_LENGTH = 40  
  
STRING_DTYPE = '<U12'  
  
PADDING_TOKEN = 'PAD'  
  
LIMIT_OF_TOKENS = 100000
```

In [0]:

```
from nltk import word_tokenize  
import numpy as np  
import string  
import re  
  
def tokenize_string(_string):  
    return [tok_.lower() for tok_ in word_tokenize(_string) if not re.fullmatch('[' + stri  
  
def pad(A, length):  
    arr = np.empty(length, dtype = STRING_DTYPE)  
    arr.fill(PADDING_TOKEN)  
    arr[:len(A)] = A  
    return arr  
  
def tokenize_row(_sentence):  
    return pad(tokenize_string(_sentence))[:MAX_LENGTH], MAX_LENGTH  
  
def encode_row(_label):  
    return 1 if _label == 'positive' else 0  
  
def encode_and_tokenize(_dataframe):  
  
    tttt = _dataframe.apply(lambda row: tokenize_row(row['review']), axis = 1)  
    llll = _dataframe.apply(lambda row: encode_row(row['sentiment']), axis = 1)  
  
    data_dict_ = { 'label': llll, 'tokens': tttt }  
  
    encoded_and_tokenized_ = pd.DataFrame(data_dict_, columns = ['label', 'tokens'])  
  
    return encoded_and_tokenized_
```

In [0]:

```
df_train_tokenized = encode_and_tokenize(df_train)
df_test_tokenized = encode_and_tokenize(df_test)
```

In [0]:

```
from collections import Counter

def get_tokens_list(_dataframe):
    all_tokens_ = []
    for sent_ in _dataframe['tokens'].values:
        all_tokens_.extend(sent_)
    tokens_counter_ = Counter(all_tokens_)
    return [t for t, _ in tokens_counter_.most_common(LIMIT_OF_TOKENS)]
```

In [0]:

```
tokens_list = get_tokens_list(pd.concat([df_train_tokenized, df_test_tokenized]))
```

In [0]:

```
word_to_int_dict = {}
word_to_int_dict.update(
    {t : i for i, t in enumerate(tokens_list)})
```

In [0]:

```
def intize_row(_tokens):
    return np.array([word_to_int_dict[t]
                    if t in word_to_int_dict
                    else 0
                    for t in _tokens])

def encode_and_tokenize(_dataframe):
    iii = _dataframe.apply(lambda row: intize_row(row['tokens']), axis = 1)
    data_dict_ = { 'label': _dataframe['label'], 'ints': iii }
    intized_ = pd.DataFrame(data_dict_, columns = ['label', 'ints'])
    return intized_
```

In [0]:

```
df_train_intized = encode_and_tokenize(df_train_tokenized)
df_test_intized = encode_and_tokenize(df_test_tokenized)
```

## Задание 2

Реализуйте и обучите двунаправленную рекуррентную сеть (*LSTM* или *GRU*).

Какого качества классификации удалось достичь?

In [18]:

```
! pip install tensorflow-gpu --pre --quiet
```

```
! pip show tensorflow-gpu
```

```
|████████████████████| 516.2MB 30kB/s
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: keras-preprocessing, wrapt, gast, grpcio, astunparse, h5py, termcolor, tensorflow-estimator, six, google-pasta, absl-py, opt-einsum, wheel, numpy, protobuf, scipy, tensorboard
Required-by:
```

In [0]:

```
import tensorflow as tf
from tensorflow import keras
```

In [0]:

```
# To fix memory leak: https://github.com/tensorflow/tensorflow/issues/33009
tf.compat.v1.disable_eager_execution()
```

Здесь будем использовать такую конфигурацию рекуррентного *LSTM*-слоя, которая позволит использовать очень быструю *cuDNN* имплементацию.

In [21]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Bidirectional, LSTM, Dense

# The requirements to use the cuDNN implementation are:
# 1. `activation` == `tanh`
# 2. `recurrent_activation` == `sigmoid`
# 3. `recurrent_dropout` == 0
# 4. `unroll` is `False`
# 5. `use_bias` is `True`
# 6. `reset_after` is `True`
# 7. Inputs, if use masking, are strictly right-padded.

model = tf.keras.Sequential()

model.add(Bidirectional(LSTM(100, return_sequences = False), merge_mode = 'concat',
           input_shape = (MAX_LENGTH, 1)))
model.add(Dense(1, activation = 'sigmoid'))
```

WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cudnn kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cudnn kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer lstm will not use cuDNN kernel since it doesn't meet the cudnn kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/resource\_variable\_ops.py:1666: calling BaseResourceVariable.\_\_init\_\_(from tensorflow.python.ops.resource\_variable\_ops) with constraint is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass \*\_constraint arguments to layers.

In [22]:

```
model.compile(optimizer = 'adam',
              loss = 'binary_crossentropy',
              metrics = ['accuracy'])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
bidirectional (Bidirectional (None, 200))		81600
=====		
dense (Dense)	(None, 1)	201
=====		
Total params: 81,801		
Trainable params: 81,801		
Non-trainable params: 0		

In [0]:

```
X_train_intized = np.asarray(list(df_train_intized['ints'].values), dtype = float)[..., np.newaxis]
X_test_intized = np.asarray(list(df_test_intized['ints'].values), dtype = float)[..., np.newaxis]

y_train_intized = np.asarray(list(df_train_intized['label'].values))
y_test_intized = np.asarray(list(df_test_intized['label'].values))
```

In [24]:

```
model.fit(x = X_train_intized, y = y_train_intized, validation_split = 0.15, epochs = 20)
```

Train on 38250 samples, validate on 6750 samples  
Epoch 1/20  
38250/38250 [=====] - 51s 1ms/sample - loss: 0.6936  
- accuracy: 0.5191 - val\_loss: 0.6901 - val\_accuracy: 0.5313  
Epoch 2/20  
38250/38250 [=====] - 51s 1ms/sample - loss: 0.6884  
- accuracy: 0.5418 - val\_loss: 0.6862 - val\_accuracy: 0.5560  
Epoch 3/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6853  
- accuracy: 0.5501 - val\_loss: 0.6826 - val\_accuracy: 0.5612  
Epoch 4/20  
38250/38250 [=====] - 51s 1ms/sample - loss: 0.6828  
- accuracy: 0.5585 - val\_loss: 0.6889 - val\_accuracy: 0.5498  
Epoch 5/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6808  
- accuracy: 0.5618 - val\_loss: 0.6818 - val\_accuracy: 0.5630  
Epoch 6/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6787  
- accuracy: 0.5702 - val\_loss: 0.6788 - val\_accuracy: 0.5695  
Epoch 7/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6768  
- accuracy: 0.5738 - val\_loss: 0.6825 - val\_accuracy: 0.5637  
Epoch 8/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6742  
- accuracy: 0.5739 - val\_loss: 0.6775 - val\_accuracy: 0.5720  
Epoch 9/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6724  
- accuracy: 0.5778 - val\_loss: 0.6807 - val\_accuracy: 0.5615  
Epoch 10/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6701  
- accuracy: 0.5816 - val\_loss: 0.6777 - val\_accuracy: 0.5693  
Epoch 11/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6676  
- accuracy: 0.5858 - val\_loss: 0.6787 - val\_accuracy: 0.5637  
Epoch 12/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6654  
- accuracy: 0.5880 - val\_loss: 0.6768 - val\_accuracy: 0.5680  
Epoch 13/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6621  
- accuracy: 0.5933 - val\_loss: 0.6726 - val\_accuracy: 0.5766  
Epoch 14/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6598  
- accuracy: 0.5970 - val\_loss: 0.6755 - val\_accuracy: 0.5701  
Epoch 15/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6571  
- accuracy: 0.5987 - val\_loss: 0.6728 - val\_accuracy: 0.5692  
Epoch 16/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6542  
- accuracy: 0.6018 - val\_loss: 0.6747 - val\_accuracy: 0.5661  
Epoch 17/20  
38250/38250 [=====] - 50s 1ms/sample - loss: 0.6516  
- accuracy: 0.6045 - val\_loss: 0.6726 - val\_accuracy: 0.5736  
Epoch 18/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6485  
- accuracy: 0.6095 - val\_loss: 0.6807 - val\_accuracy: 0.5759  
Epoch 19/20  
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6467

```
- accuracy: 0.6125 - val_loss: 0.6771 - val_accuracy: 0.5714
Epoch 20/20
38250/38250 [=====] - 49s 1ms/sample - loss: 0.6438
- accuracy: 0.6130 - val_loss: 0.6781 - val_accuracy: 0.5658
```

Out[24]:

```
<tensorflow.python.keras.callbacks.History at 0x7f8fd8a8b208>
```

In [25]:

```
results = model.evaluate(X_test_intized, y_test_intized)
print('Test loss, test accuracy:', results)
```

```
Test loss, test accuracy: [0.6761944528579712, 0.5762]
```

На валидационной выборке удалось достичь точности 57%.

## Задание 3

Используйте индексы слов и их различное внутреннее представление (*word2vec, glove*). Как влияет данное преобразование на качество классификации?

Используем 300-мерные вектора *FastText* — лучшую на сегодняшний день имплементацию word2vec: <https://fasttext.cc/docs/en/english-vectors.html> (<https://fasttext.cc/docs/en/english-vectors.html>). Файл пришлось доработать — 9-я строка не читалась.

In [0]:

```
# VECTORS_ARCHIVE_NAME = 'wiki-news-300d-1M-fixed.zip'
# VECTORS_FILE_NAME = 'wiki-news-300d-1M-fixed.vec'
# VECTORS_LOCAL_DIR_NAME = 'vectors'
```

In [0]:

```
# with ZipFile(os.path.join(BASE_DIR, VECTORS_ARCHIVE_NAME), 'r') as zip_:
#     zip_.extractall(VECTORS_LOCAL_DIR_NAME)
```

Создадим уменьшенный словарь, содержащий только встреченные токены, чтобы уменьшить нагрузку на Google Drive:

In [0]:

```
# def build_vectors_dict(_actual_tokens, _vectors_file_path, _unknown_token = 'unknown'):
#     vec_data_ = pd.read_csv(_vectors_file_path, sep = ' ', header = None, skiprows = [9])
#     actual_vectors_ = [x for x in vec_data_.values if x[0] in _actual_tokens or x[0] == _unknown_token]
#     return actual_vectors_
```

In [0]:

```
# actual_vectors = build_vectors_dict(tokens_list, os.path.join(VECTORS_LOCAL_DIR_NAME, VEC
```

In [0]:

```
# vectors_np = np.array(actual_vectors)

# vectors_dict = dict(zip(vectors_np[:, 0], vectors_np[:, 1:]))

# vectors_dict_file_name = 'word-vec-dict-{}-items'.format(len(vectors_dict))

# vectors_dict_file_path = os.path.join(BASE_DIR, vectors_dict_file_name)

# np.savez_compressed(vectors_dict_file_path, vectors_dict, allow_pickle = True)
```

In [0]:

```
vectors_dict_file_path = './drive/My Drive/Colab Files/mo-2/word-vec-dict-56485-items.npz'
```

In [0]:

```
vectors_dict_data = np.load(vectors_dict_file_path, allow_pickle = True)

vectors_dict = vectors_dict_data['arr_0'][()]
```

In [0]:

```
VECTORS_LENGTH = 300
```

In [0]:

```
def tokens_to_vectors(_word_to_vec_dict, _tokens, _unknown_token):
    return [_word_to_vec_dict[t]
            if t in _word_to_vec_dict
            else _word_to_vec_dict[_unknown_token]
            for t in _tokens]

def row_to_vectors(_tokens):
    return np.array(tokens_to_vectors(vectors_dict, _tokens, 'unknown'))

def vectorize(_dataframe):

    vvvv = _dataframe.apply(lambda row: row_to_vectors(row['tokens']), axis = 1)

    data_dict_ = { 'label': _dataframe['label'], 'vectors': vvvv }

    vectorized_ = pd.DataFrame(data_dict_, columns = ['label', 'vectors'])

    return vectorized_
```

In [0]:

```
df_train_vectorized = vectorize(df_train_tokenized)
df_test_vectorized = vectorize(df_test_tokenized)
```

In [0]:

```
X_train_vectorized = np.asarray(list(df_train_vectorized['vectors'].values), dtype = float)
X_test_vectorized = np.asarray(list(df_test_vectorized['vectors'].values), dtype = float)

y_train_vectorized = np.asarray(list(df_train_vectorized['label'].values))
y_test_vectorized = np.asarray(list(df_test_vectorized['label'].values))
```

In [37]:

```
model_2 = tf.keras.Sequential()

model_2.add(Bidirectional(LSTM(100, return_sequences = False), merge_mode = 'concat',
                         input_shape = (MAX_LENGTH, VECTORS_LENGTH)))
model_2.add(Dense(1, activation = 'sigmoid'))
```

WARNING:tensorflow:Layer lstm\_1 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer lstm\_1 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer lstm\_1 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

In [38]:

```
model_2.compile(optimizer = 'adam',
                 loss = 'binary_crossentropy',
                 metrics = ['accuracy'])
```

```
model_2.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
bidirectional_1 (Bidirection (None, 200)		320800
dense_1 (Dense)	(None, 1)	201
Total params:	321,001	
Trainable params:	321,001	
Non-trainable params:	0	

In [39]:

```
model_2.fit(x = X_train_vectorized, y = y_train_vectorized, validation_split = 0.15, epochs=20)

Train on 38250 samples, validate on 6750 samples
Epoch 1/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.5456
- accuracy: 0.7143 - val_loss: 0.4991 - val_accuracy: 0.7511
Epoch 2/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.4851
- accuracy: 0.7566 - val_loss: 0.4786 - val_accuracy: 0.7603
Epoch 3/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.4589
- accuracy: 0.7734 - val_loss: 0.4749 - val_accuracy: 0.7640
Epoch 4/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.4354
- accuracy: 0.7882 - val_loss: 0.4676 - val_accuracy: 0.7686
Epoch 5/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.4117
- accuracy: 0.8025 - val_loss: 0.4917 - val_accuracy: 0.7613
Epoch 6/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.3849
- accuracy: 0.8179 - val_loss: 0.4727 - val_accuracy: 0.7667
Epoch 7/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.3511
- accuracy: 0.8372 - val_loss: 0.5012 - val_accuracy: 0.7529
Epoch 8/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.3122
- accuracy: 0.8599 - val_loss: 0.5162 - val_accuracy: 0.7553
Epoch 9/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.2662
- accuracy: 0.8848 - val_loss: 0.5886 - val_accuracy: 0.7609
Epoch 10/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.2165
- accuracy: 0.9092 - val_loss: 0.6468 - val_accuracy: 0.7603
Epoch 11/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.1680
- accuracy: 0.9329 - val_loss: 0.7128 - val_accuracy: 0.7470
Epoch 12/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.1269
- accuracy: 0.9524 - val_loss: 0.8222 - val_accuracy: 0.7526
Epoch 13/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0908
- accuracy: 0.9670 - val_loss: 0.8999 - val_accuracy: 0.7487
Epoch 14/20
38250/38250 [=====] - 53s 1ms/sample - loss: 0.0651
- accuracy: 0.9780 - val_loss: 1.0383 - val_accuracy: 0.7427
Epoch 15/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0495
- accuracy: 0.9842 - val_loss: 1.1222 - val_accuracy: 0.7464
Epoch 16/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0379
- accuracy: 0.9884 - val_loss: 1.2738 - val_accuracy: 0.7513
Epoch 17/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0326
- accuracy: 0.9897 - val_loss: 1.3313 - val_accuracy: 0.7434
Epoch 18/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0364
- accuracy: 0.9888 - val_loss: 1.2889 - val_accuracy: 0.7495
Epoch 19/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0237
```

```
- accuracy: 0.9933 - val_loss: 1.4557 - val_accuracy: 0.7465
Epoch 20/20
38250/38250 [=====] - 52s 1ms/sample - loss: 0.0216
- accuracy: 0.9940 - val_loss: 1.3947 - val_accuracy: 0.7459
```

Out[39]:

```
<tensorflow.python.keras.callbacks.History at 0x7f8fd4befa20>
```

In [40]:

```
results_2 = model_2.evaluate(X_test_vectorized, y_test_vectorized)

print('Test loss, test accuracy:', results_2)
```

```
Test loss, test accuracy: [1.3811187601089479, 0.7412]
```

Как и ожидалось, использование эмбеддингов показало лучший результат, чем кодирование слов просто целыми числами — 74%.

## Задание 4

Поэкспериментируйте со структурой сети (добавьте больше рекуррентных, полносвязных или сверточных слоев). Как это повлияло на качество классификации?

In [41]:

```
model_3 = tf.keras.Sequential()

model_3.add(Bidirectional(LSTM(5, return_sequences = True), merge_mode = 'concat',
                         input_shape = (MAX_LENGTH, VECTORS_LENGTH)))
model_3.add(LSTM(1, return_sequences = False))
model_3.add(Dense(10, activation = 'linear'))
model_3.add(Dense(1, activation = 'sigmoid'))
```

```
WARNING:tensorflow:Layer lstm_2 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU
```

```
WARNING:tensorflow:Layer lstm_2 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU
```

```
WARNING:tensorflow:Layer lstm_2 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU
```

```
WARNING:tensorflow:Layer lstm_3 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU
```

In [42]:

```
model_3.compile(optimizer = 'adam',
                 loss = 'binary_crossentropy',
                 metrics = ['accuracy'])
```

```
model_3.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
bidirectional_2 (Bidirection (None, 40, 10)		12240
lstm_3 (LSTM)	(None, 1)	48
dense_2 (Dense)	(None, 10)	20
dense_3 (Dense)	(None, 1)	11
=====		
Total params:	12,319	
Trainable params:	12,319	
Non-trainable params:	0	

In [44]:

```
model_3.fit(x = X_train_vectorized, y = y_train_vectorized, validation_split = 0.15, epochs=20)

Train on 38250 samples, validate on 6750 samples
Epoch 1/20
38250/38250 [=====] - 83s 2ms/sample - loss: 0.6033
- accuracy: 0.6808 - val_loss: 0.5700 - val_accuracy: 0.7108
Epoch 2/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.5302
- accuracy: 0.7388 - val_loss: 0.5133 - val_accuracy: 0.7453
Epoch 3/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4987
- accuracy: 0.7546 - val_loss: 0.4906 - val_accuracy: 0.7566
Epoch 4/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4819
- accuracy: 0.7621 - val_loss: 0.4779 - val_accuracy: 0.7625
Epoch 5/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.4681
- accuracy: 0.7725 - val_loss: 0.4778 - val_accuracy: 0.7579
Epoch 6/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4602
- accuracy: 0.7770 - val_loss: 0.4755 - val_accuracy: 0.7636
Epoch 7/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4515
- accuracy: 0.7817 - val_loss: 0.4838 - val_accuracy: 0.7563
Epoch 8/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4444
- accuracy: 0.7843 - val_loss: 0.4796 - val_accuracy: 0.7613
Epoch 9/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.4363
- accuracy: 0.7929 - val_loss: 0.4616 - val_accuracy: 0.7683
Epoch 10/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4329
- accuracy: 0.7923 - val_loss: 0.4661 - val_accuracy: 0.7726
Epoch 11/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4258
- accuracy: 0.7961 - val_loss: 0.4658 - val_accuracy: 0.7674
Epoch 12/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4192
- accuracy: 0.8000 - val_loss: 0.4680 - val_accuracy: 0.7711
Epoch 13/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.4138
- accuracy: 0.8031 - val_loss: 0.4609 - val_accuracy: 0.7754
Epoch 14/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.4084
- accuracy: 0.8092 - val_loss: 0.4622 - val_accuracy: 0.7763
Epoch 15/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.4028
- accuracy: 0.8098 - val_loss: 0.4717 - val_accuracy: 0.7705
Epoch 16/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.3975
- accuracy: 0.8142 - val_loss: 0.4648 - val_accuracy: 0.7742
Epoch 17/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.3942
- accuracy: 0.8166 - val_loss: 0.4751 - val_accuracy: 0.7673
Epoch 18/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.3886
- accuracy: 0.8202 - val_loss: 0.4730 - val_accuracy: 0.7730
Epoch 19/20
38250/38250 [=====] - 81s 2ms/sample - loss: 0.3844
```

```
- accuracy: 0.8226 - val_loss: 0.4740 - val_accuracy: 0.7753
Epoch 20/20
38250/38250 [=====] - 82s 2ms/sample - loss: 0.3797
- accuracy: 0.8246 - val_loss: 0.4765 - val_accuracy: 0.7727
```

Out[44]:

```
<tensorflow.python.keras.callbacks.History at 0x7f8ef6e7a6a0>
```

In [45]:

```
results_3 = model_3.evaluate(X_test_vectorized, y_test_vectorized)

print('Test loss, test accuracy:', results_3)
```

```
Test loss, test accuracy: [0.49010655212402343, 0.766]
```

Добавление ещё одного рекуррентного слоя ненамного улучшило результат — точность 76% на тестовой выборке.

## Задание 5

Используйте предобученную рекуррентную нейронную сеть (например, *DeepMoji* или что-то подобное).

Какой максимальный результат удалось получить на контрольной выборке?

На своих моделях удалось достичнуть максимальной точности 76%.

# Лабораторная работа №8

## Рекуррентные нейронные сети для анализа временных рядов

Набор данных для прогнозирования временных рядов, который состоит из среднемесячного числа пятен на солнце, наблюдавшихся с января 1749 по август 2017.

Данные в виде csv-файла можно скачать на сайте Kaggle: <https://www.kaggle.com/robervalt/sunspots/> (<https://www.kaggle.com/robervalt/sunspots/>)

### Задание 1

Загрузите данные. Изобразите ряд в виде графика. Вычислите основные характеристики временного ряда (сезонность, тренд, автокорреляцию).

In [1]:

```
from google.colab import drive  
drive.mount('/content/drive', force_remount = True)
```

Mounted at /content/drive

In [0]:

```
BASE_DIR = '/content/drive/My Drive/Colab Files/mo-2'  
  
import sys  
  
sys.path.append(BASE_DIR)  
  
import os
```

In [0]:

```
DATA_ARCHIVE_NAME = 'sunspots.zip'  
  
LOCAL_DIR_NAME = 'sunspots'
```

In [0]:

```
from zipfile import ZipFile  
  
with ZipFile(os.path.join(BASE_DIR, DATA_ARCHIVE_NAME), 'r') as zip_:  
    zip_.extractall(LOCAL_DIR_NAME)
```

In [0]:

```
DATA_FILE_PATH = 'sunspots/Sunspots.csv'
```

In [0]:

```
import pandas as pd

all_df = pd.read_csv(DATA_FILE_PATH, parse_dates = ['Date'], index_col = 'Date')
```

In [7]:

```
print(all_df.shape)

(3252, 2)
```

In [8]:

```
all_df.keys()
```

Out[8]:

```
Index(['Unnamed: 0', 'Monthly Mean Total Sunspot Number'], dtype='object')
```

In [9]:

```
from statsmodels.tsa.seasonal import seasonal_decompose

additive = seasonal_decompose(all_df['Monthly Mean Total Sunspot Number'], model = 'additive')

/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated. Use the functions in the public API at pandas.testing instead.
    import pandas.util.testing as tm
```

In [0]:

```
%matplotlib inline

import matplotlib.pyplot as plt
```

In [0]:

```
import seaborn as sns

from matplotlib import rcParams

rcParams['figure.figsize'] = 11.7, 8.27

sns.set()

sns.set_palette(sns.color_palette('hls', 8))
```

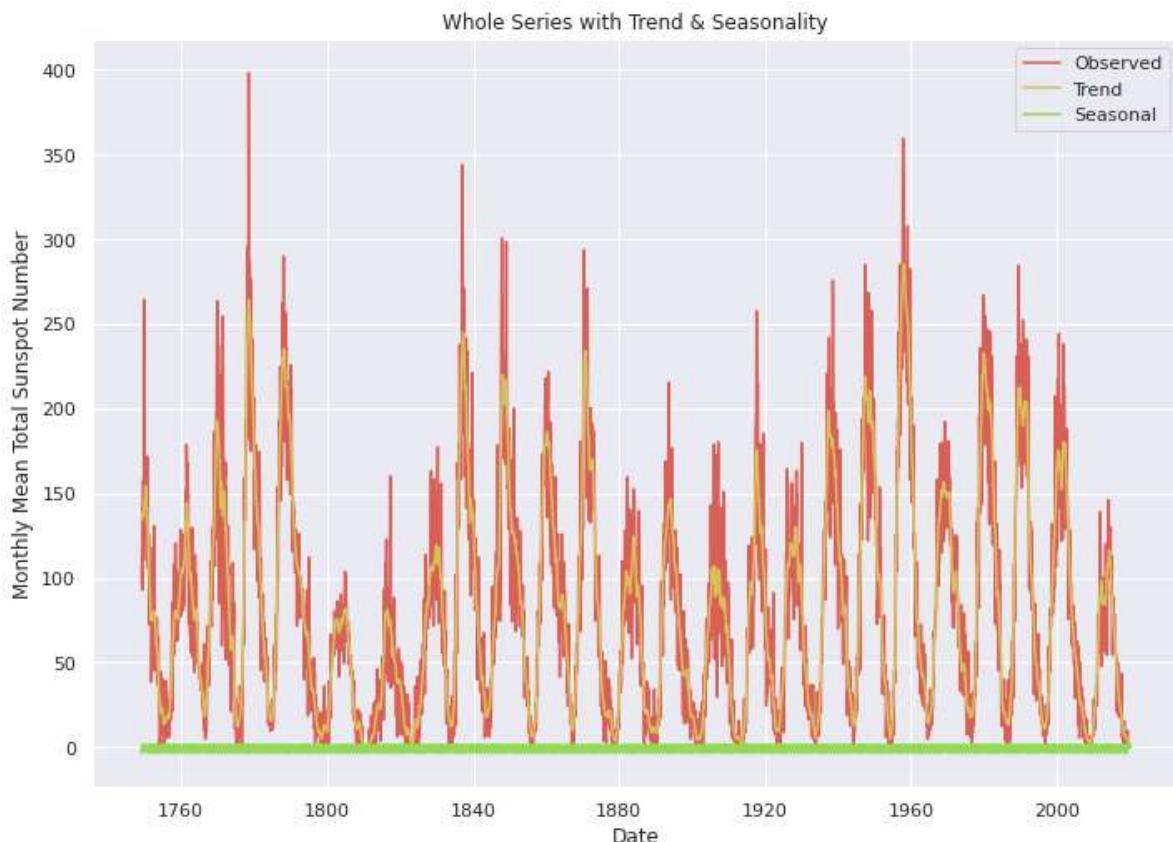
In [12]:

```
sns.lineplot(data = additive.observed, label = 'Observed')
sns.lineplot(data = additive.trend, label = 'Trend')
sns.lineplot(data = additive.seasonal, label = 'Seasonal')

plt.xlabel('Date')
plt.ylabel('Monthly Mean Total Sunspot Number')

plt.title('Whole Series with Trend & Seasonality')

plt.show()
```



Рассмотрим подробнее на небольшом промежутке:

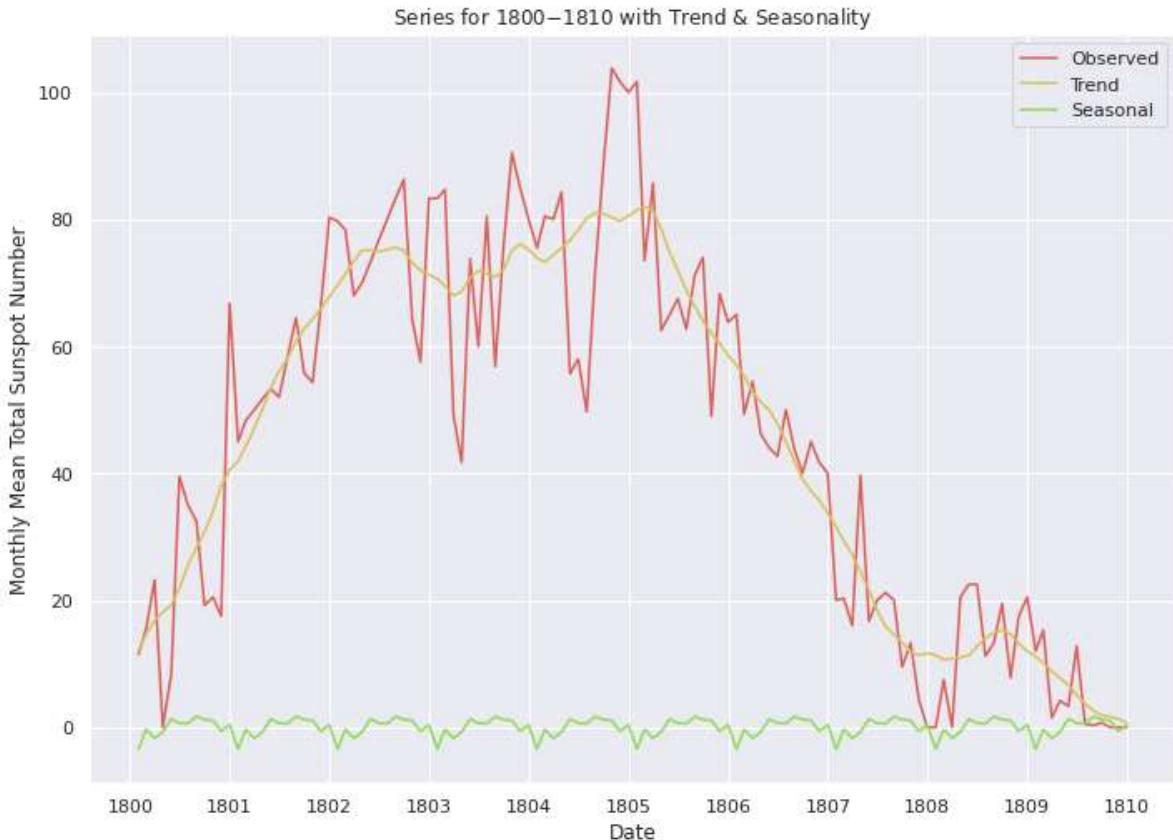
In [13]:

```
sns.lineplot(data = additive.observed['1800-01-01':'1810-01-01'], label = 'Observed')
sns.lineplot(data = additive.trend['1800-01-01':'1810-01-01'], label = 'Trend')
sns.lineplot(data = additive.seasonal['1800-01-01':'1810-01-01'], label = 'Seasonal')

plt.xlabel('Date')
plt.ylabel('Monthly Mean Total Sunspot Number')

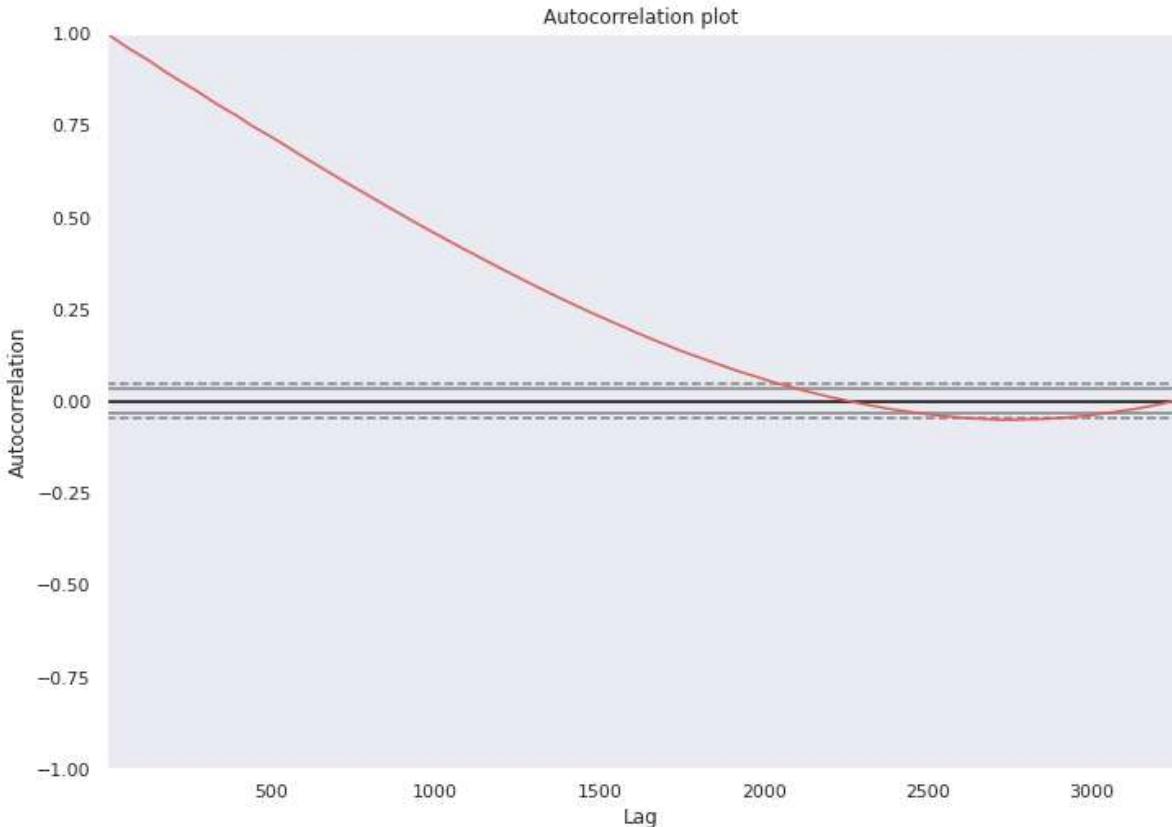
plt.title('Series for 1800$-$1810 with Trend & Seasonality')

plt.show()
```



In [14]:

```
from pandas.plotting import autocorrelation_plot  
  
autocorrelation_plot(all_df.values.tolist())  
  
plt.title('Autocorrelation plot')  
  
plt.show()
```



## Задание 2

Для прогнозирования разделите временной ряд на обучающую, валидационную и контрольную выборки.

Этот шаг будет применён автоматически как параметр `validation_split` метода `model.fit()`.

## Задание 3

Примените модель ARIMA для прогнозирования значений данного временного ряда.

In [15]:

```
! pip install pmdarima  
!  
! pip show pmdarima
```

```
Requirement already satisfied: pmdarima in /usr/local/lib/python3.6/dist-packages (1.5.3)  
Requirement already satisfied: Cython>=0.29 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (0.29.16)  
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (1.0.3)  
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (0.22.2.post1)  
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (1.18.2)  
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (1.4.1)  
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (0.14.1)  
Requirement already satisfied: statsmodels>=0.10.2 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (0.10.2)  
Requirement already satisfied: urllib3 in /usr/local/lib/python3.6/dist-packages (from pmdarima) (1.24.3)  
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19->pmdarima) (2018.9)  
Requirement already satisfied: python-dateutil>=2.6.1 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19->pmdarima) (2.8.1)  
Requirement already satisfied: patsy>=0.4.0 in /usr/local/lib/python3.6/dist-packages (from statsmodels>=0.10.2->pmdarima) (0.5.1)  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-packages (from python-dateutil>=2.6.1->pandas>=0.19->pmdarima) (1.12.0)  
Name: pmdarima  
Version: 1.5.3  
Summary: Python's forecast::auto.arima equivalent  
Home-page: http://alkaline-ml.com/pmdarima (http://alkaline-ml.com/pmdarima)  
Author: None  
Author-email: None  
License: MIT  
Location: /usr/local/lib/python3.6/dist-packages  
Requires: joblib, Cython, pandas, urllib3, statsmodels, numpy, scikit-learn, scipy  
Required-by:
```

In [0]:

```
test_period = 6 * 12
```

In [17]:

```
from pmdarima.arima import auto_arima

arima = auto_arima(all_df[ 'Monthly Mean Total Sunspot Number'][:-test_period],
                  trace = True, error_action = 'ignore', suppress_warnings = True, seasonal=
```

Performing stepwise search to minimize aic  
Fit ARIMA: (2, 0, 2)x(1, 0, 1, 12) (constant=True); AIC=29588.066, BIC=2963  
6.583, Time=23.372 seconds  
Fit ARIMA: (0, 0, 0)x(0, 0, 0, 12) (constant=True); AIC=35872.649, BIC=3588  
4.778, Time=0.065 seconds  
Fit ARIMA: (1, 0, 0)x(1, 0, 0, 12) (constant=True); AIC=30025.545, BIC=3004  
9.804, Time=5.185 seconds  
Fit ARIMA: (0, 0, 1)x(0, 0, 1, 12) (constant=True); AIC=32380.500, BIC=3240  
4.758, Time=5.341 seconds  
Fit ARIMA: (0, 0, 0)x(0, 0, 0, 12) (constant=False); AIC=38764.997, BIC=3877  
1.062, Time=0.049 seconds  
Fit ARIMA: (2, 0, 2)x(0, 0, 1, 12) (constant=True); AIC=29586.807, BIC=2962  
9.260, Time=13.302 seconds  
Fit ARIMA: (2, 0, 2)x(0, 0, 0, 12) (constant=True); AIC=29595.958, BIC=2963  
2.345, Time=2.032 seconds  
Fit ARIMA: (2, 0, 2)x(0, 0, 2, 12) (constant=True); AIC=29587.374, BIC=2963  
5.891, Time=45.942 seconds  
Fit ARIMA: (2, 0, 2)x(1, 0, 0, 12) (constant=True); AIC=29587.365, BIC=2962  
9.817, Time=16.549 seconds  
Fit ARIMA: (2, 0, 2)x(1, 0, 2, 12) (constant=True); AIC=29582.451, BIC=2963  
7.033, Time=67.935 seconds  
Fit ARIMA: (2, 0, 2)x(2, 0, 2, 12) (constant=True); AIC=29552.494, BIC=2961  
3.140, Time=71.844 seconds  
Near non-invertible roots for order (2, 0, 2)(2, 0, 2, 12); setting score to  
inf (at least one inverse root too close to the border of the unit circle:  
0.993)  
Fit ARIMA: (2, 0, 2)x(2, 0, 1, 12) (constant=True); AIC=29581.341, BIC=2963  
5.923, Time=64.230 seconds  
Fit ARIMA: (2, 0, 2)x(2, 0, 0, 12) (constant=True); AIC=29587.014, BIC=2963  
5.531, Time=52.283 seconds  
Fit ARIMA: (1, 0, 2)x(2, 0, 1, 12) (constant=True); AIC=29579.106, BIC=2962  
7.623, Time=43.009 seconds  
Fit ARIMA: (1, 0, 2)x(1, 0, 1, 12) (constant=True); AIC=29586.021, BIC=2962  
8.473, Time=18.368 seconds  
Fit ARIMA: (1, 0, 2)x(2, 0, 0, 12) (constant=True); AIC=29585.061, BIC=2962  
7.513, Time=39.560 seconds  
Fit ARIMA: (1, 0, 2)x(2, 0, 2, 12) (constant=True); AIC=29551.544, BIC=2960  
6.126, Time=72.061 seconds  
Near non-invertible roots for order (1, 0, 2)(2, 0, 2, 12); setting score to  
inf (at least one inverse root too close to the border of the unit circle:  
0.995)  
Fit ARIMA: (1, 0, 2)x(1, 0, 0, 12) (constant=True); AIC=29585.380, BIC=2962  
1.768, Time=8.778 seconds  
Fit ARIMA: (1, 0, 2)x(1, 0, 2, 12) (constant=True); AIC=29580.389, BIC=2962  
8.906, Time=49.857 seconds  
Fit ARIMA: (0, 0, 2)x(2, 0, 1, 12) (constant=True); AIC=31382.620, BIC=3142  
5.073, Time=33.173 seconds  
Near non-invertible roots for order (0, 0, 2)(2, 0, 1, 12); setting score to  
inf (at least one inverse root too close to the border of the unit circle:  
1.000)  
Fit ARIMA: (1, 0, 1)x(2, 0, 1, 12) (constant=True); AIC=29629.733, BIC=2967  
2.185, Time=54.417 seconds  
Fit ARIMA: (1, 0, 3)x(2, 0, 1, 12) (constant=True); AIC=29580.861, BIC=2963  
5.443, Time=48.949 seconds

Fit ARIMA: (0, 0, 1)x(2, 0, 1, 12) (constant=True); AIC=32043.176, BIC=32079.563, Time=28.353 seconds  
Near non-invertible roots for order (0, 0, 1)(2, 0, 1, 12); setting score to inf (at least one inverse root too close to the border of the unit circle:  
1.000)  
Fit ARIMA: (0, 0, 3)x(2, 0, 1, 12) (constant=True); AIC=31074.278, BIC=31122.795, Time=42.232 seconds  
Near non-invertible roots for order (0, 0, 3)(2, 0, 1, 12); setting score to inf (at least one inverse root too close to the border of the unit circle:  
1.000)  
Fit ARIMA: (2, 0, 1)x(2, 0, 1, 12) (constant=True); AIC=29592.648, BIC=29641.165, Time=51.681 seconds  
Fit ARIMA: (2, 0, 3)x(2, 0, 1, 12) (constant=True); AIC=29579.013, BIC=29639.659, Time=62.712 seconds  
Fit ARIMA: (2, 0, 3)x(1, 0, 1, 12) (constant=True); AIC=29586.998, BIC=29641.579, Time=22.034 seconds  
Fit ARIMA: (2, 0, 3)x(2, 0, 0, 12) (constant=True); AIC=29585.838, BIC=29640.420, Time=60.607 seconds  
Fit ARIMA: (2, 0, 3)x(2, 0, 2, 12) (constant=True); AIC=29543.930, BIC=29610.641, Time=80.082 seconds  
Near non-invertible roots for order (2, 0, 3)(2, 0, 2, 12); setting score to inf (at least one inverse root too close to the border of the unit circle:  
0.995)  
Fit ARIMA: (2, 0, 3)x(1, 0, 0, 12) (constant=True); AIC=29586.591, BIC=29635.108, Time=17.100 seconds  
Fit ARIMA: (2, 0, 3)x(1, 0, 2, 12) (constant=True); AIC=29580.547, BIC=29641.194, Time=72.056 seconds  
Fit ARIMA: (3, 0, 3)x(2, 0, 1, 12) (constant=True); AIC=29578.087, BIC=29644.798, Time=79.448 seconds  
Fit ARIMA: (3, 0, 3)x(1, 0, 1, 12) (constant=True); AIC=29586.766, BIC=29647.413, Time=22.877 seconds  
Fit ARIMA: (3, 0, 3)x(2, 0, 0, 12) (constant=True); AIC=29584.745, BIC=29645.391, Time=79.798 seconds  
Fit ARIMA: (3, 0, 3)x(2, 0, 2, 12) (constant=True); AIC=29546.099, BIC=29618.874, Time=82.080 seconds  
Near non-invertible roots for order (3, 0, 3)(2, 0, 2, 12); setting score to inf (at least one inverse root too close to the border of the unit circle:  
0.995)  
Fit ARIMA: (3, 0, 3)x(1, 0, 0, 12) (constant=True); AIC=29585.293, BIC=29639.875, Time=25.847 seconds  
Fit ARIMA: (3, 0, 3)x(1, 0, 2, 12) (constant=True); AIC=29579.288, BIC=29645.999, Time=71.712 seconds  
Fit ARIMA: (3, 0, 2)x(2, 0, 1, 12) (constant=True); AIC=29587.135, BIC=29647.781, Time=73.021 seconds  
Fit ARIMA: (4, 0, 3)x(2, 0, 1, 12) (constant=True); AIC=29579.706, BIC=29652.482, Time=82.381 seconds  
Fit ARIMA: (3, 0, 4)x(2, 0, 1, 12) (constant=True); AIC=29579.646, BIC=29652.421, Time=91.857 seconds  
Fit ARIMA: (2, 0, 4)x(2, 0, 1, 12) (constant=True); AIC=29576.620, BIC=29643.331, Time=70.104 seconds  
Fit ARIMA: (2, 0, 4)x(1, 0, 1, 12) (constant=True); AIC=29583.960, BIC=29644.607, Time=29.800 seconds  
Fit ARIMA: (2, 0, 4)x(2, 0, 0, 12) (constant=True); AIC=29582.698, BIC=29643.345, Time=71.533 seconds  
Fit ARIMA: (2, 0, 4)x(2, 0, 2, 12) (constant=True); AIC=29547.505, BIC=29620.281, Time=100.620 seconds  
Near non-invertible roots for order (2, 0, 4)(2, 0, 2, 12); setting score to inf (at least one inverse root too close to the border of the unit circle:  
0.994)  
Fit ARIMA: (2, 0, 4)x(1, 0, 0, 12) (constant=True); AIC=29568.222, BIC=29622.804, Time=23.695 seconds

```
Fit ARIMA: (2, 0, 4)x(0, 0, 0, 12) (constant=True); AIC=29479.053, BIC=2952  
7.570, Time=8.133 seconds  
Fit ARIMA: (2, 0, 4)x(0, 0, 1, 12) (constant=True); AIC=29472.268, BIC=2952  
6.849, Time=31.449 seconds  
Fit ARIMA: (2, 0, 4)x(0, 0, 2, 12) (constant=True); AIC=29573.231, BIC=2963  
3.878, Time=81.292 seconds  
Fit ARIMA: (2, 0, 4)x(1, 0, 2, 12) (constant=True); AIC=29578.088, BIC=2964  
4.799, Time=84.294 seconds  
Fit ARIMA: (1, 0, 4)x(0, 0, 1, 12) (constant=True); AIC=29579.628, BIC=2962  
8.145, Time=15.428 seconds  
Fit ARIMA: (2, 0, 3)x(0, 0, 1, 12) (constant=True); AIC=29585.945, BIC=2963  
4.462, Time=10.997 seconds  
Fit ARIMA: (3, 0, 4)x(0, 0, 1, 12) (constant=True); AIC=29523.347, BIC=2958  
3.993, Time=30.866 seconds  
Fit ARIMA: (2, 0, 5)x(0, 0, 1, 12) (constant=True); AIC=29583.600, BIC=2964  
4.247, Time=15.339 seconds  
Fit ARIMA: (1, 0, 3)x(0, 0, 1, 12) (constant=True); AIC=29586.789, BIC=2962  
9.241, Time=10.503 seconds  
Fit ARIMA: (1, 0, 5)x(0, 0, 1, 12) (constant=True); AIC=29579.796, BIC=2963  
4.378, Time=21.199 seconds  
Fit ARIMA: (3, 0, 3)x(0, 0, 1, 12) (constant=True); AIC=29584.591, BIC=2963  
9.173, Time=20.767 seconds  
Fit ARIMA: (3, 0, 5)x(0, 0, 1, 12) (constant=True); AIC=29449.139, BIC=2951  
5.850, Time=34.889 seconds  
Near non-invertible roots for order (3, 0, 5)(0, 0, 1, 12); setting score to  
inf (at least one inverse root too close to the border of the unit circle:  
0.997)  
Total fit time: 2443.158 seconds
```

In [0]:

```
arima_forecast = arima.predict(n_periods = test_period)
```

In [19]:

```
from sklearn.metrics import mean_squared_error  
  
mean_squared_error(all_df['Monthly Mean Total Sunspot Number'][~-test_period:], arima_forecast)
```

Out[19]:

```
3196.9943474969787
```

## Задание 4

Повторите эксперимент по прогнозированию, реализовав рекуррентную нейронную сеть (с как минимум 2 рекуррентными слоями).

Сначала нужно создать датасет из данных.

In [20]:

```
! pip install tensorflow-gpu --pre --quiet
```

```
! pip show tensorflow-gpu
```

```
Name: tensorflow-gpu
Version: 2.2.0rc3
Summary: TensorFlow is an open source machine learning framework for everyone.
Home-page: https://www.tensorflow.org/ (https://www.tensorflow.org/)
Author: Google Inc.
Author-email: packages@tensorflow.org
License: Apache 2.0
Location: /usr/local/lib/python3.6/dist-packages
Requires: google-pasta, opt-einsum, tensorboard, six, h5py, tensorflow-estimator, termcolor, wrapt, wheel, keras-preprocessing, protobuf, gast, astunparse, scipy, numpy, grpcio, absl-py
Required-by:
```

In [0]:

```
TIME_STEPS = 100
```

```
TEST_PERIOD = 1000
```

In [0]:

```
import numpy as np
from datetime import timezone

def timeseries_to_dataset(_X_ts, _time_steps):
    samples_n_ = len(_X_ts) - _time_steps
    print(len(_X_ts))

    X_ = np.zeros((samples_n_, _time_steps))
    y_ = np.zeros((samples_n_,))

    for i in range(samples_n_):
        X_[i] = _X_ts[i:(i + _time_steps)]
        y_[i] = _X_ts[(i + _time_steps)]

    return X_[..., np.newaxis], y_
```

In [31]:

```
X, y = timeseries_to_dataset(all_df['Monthly Mean Total Sunspot Number'][:-TEST_PERIOD].values
X_test, y_test = timeseries_to_dataset(all_df['Monthly Mean Total Sunspot Number'][-TEST_PERIOD:]
```

2252  
1000

In [0]:

```
import tensorflow as tf
from tensorflow import keras
```

In [37]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense

model = tf.keras.Sequential()

model.add(LSTM(8, activation = 'relu', return_sequences = True, input_shape = X.shape[-2:]))
model.add(LSTM(8, activation = 'relu'))
model.add(Dense(1))
```

WARNING:tensorflow:Layer lstm\_4 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

WARNING:tensorflow:Layer lstm\_5 will not use cuDNN kernel since it doesn't meet the cuDNN kernel criteria. It will use generic GPU kernel as fallback when running on GPU

In [38]:

```
model.compile(optimizer = 'adam',
              loss = 'mse',
              metrics = ['accuracy'])
```

```
model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
lstm_4 (LSTM)	(None, 100, 8)	320
lstm_5 (LSTM)	(None, 8)	544
dense_2 (Dense)	(None, 1)	9
Total params:	873	
Trainable params:	873	
Non-trainable params:	0	

In [39]:

```
model.fit(x = X, y = y, validation_split = 0.15, epochs = 20, verbose = 0)
```

Out[39]:

```
<tensorflow.python.keras.callbacks.History at 0x7f162e959fd0>
```

In [40]:

```
results = model.evaluate(X_test, y_test)

print('Test mse, test accuracy:', results)

29/29 [=====] - 1s 21ms/step - loss: 651.9083 - acc
uracy: 0.0011
Test mse, test accuracy: [651.9083251953125, 0.001111111380159855]
```

## Задание 5

Сравните качество прогноза моделей.

Какой максимальный результат удалось получить на контрольной выборке?

Нейронная сеть дала среднеквадратичную ошибку в 4 раза больше, чем ARIMA, а точность предсказания вообще близка к нулю. Можно сделать вывод, что предсказание временных рядов требует более тонкой настройки архитектуры сетей.