

# Лабораторная работа №4

## Реализация приложения по распознаванию номеров домов

Набор изображений из *Google Street View* с изображениями номеров домов, содержащий 10 классов, соответствующих цифрам от 0 до 9.

- 73257 изображений цифр в обучающей выборке;
- 26032 изображения цифр в тестовой выборке;
- 531131 изображения, которые можно использовать как дополнение к обучающей выборке;
- В двух форматах:
  - Оригинальные изображения с выделенными цифрами;
  - Изображения размером 32×32, содержащие одну цифру;
- Данные первого формата можно скачать по ссылкам:
  - <http://ufldl.stanford.edu/housenumbers/train.tar.gz> (<http://ufldl.stanford.edu/housenumbers/train.tar.gz>) (обучающая выборка);
  - <http://ufldl.stanford.edu/housenumbers/test.tar.gz> (<http://ufldl.stanford.edu/housenumbers/test.tar.gz>) (тестовая выборка);
  - <http://ufldl.stanford.edu/housenumbers/extra.tar.gz> (<http://ufldl.stanford.edu/housenumbers/extra.tar.gz>) (дополнительные данные);
- Данные второго формата можно скачать по ссылкам:
  - [http://ufldl.stanford.edu/housenumbers/train\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/train_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/train\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/train_32x32.mat)) (обучающая выборка);
  - [http://ufldl.stanford.edu/housenumbers/test\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/test_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/test\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/test_32x32.mat)) (тестовая выборка);
  - [http://ufldl.stanford.edu/housenumbers/extra\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/extra_32x32.mat) ([http://ufldl.stanford.edu/housenumbers/extra\\_32x32.mat](http://ufldl.stanford.edu/housenumbers/extra_32x32.mat)) (дополнительные данные);
- Описание данных на английском языке доступно по ссылке:
  - <http://ufldl.stanford.edu/housenumbers/> (<http://ufldl.stanford.edu/housenumbers/>)

### Задание 1

Реализуйте глубокую нейронную сеть (полносвязную или сверточную) и обучите ее на синтетических данных (например, наборы *MNIST* (<http://yann.lecun.com/exdb/mnist/>) (<http://yann.lecun.com/exdb/mnist/>) или *notMNIST*).

Ознакомьтесь с имеющимися работами по данной тематике: англоязычная статья (<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf> (<http://static.googleusercontent.com/media/research.google.com/en//pubs/archive/42241.pdf>)), видео на YouTube ([https://www.youtube.com/watch?v=vGPI\\_JvLoN0](https://www.youtube.com/watch?v=vGPI_JvLoN0)) ([https://www.youtube.com/watch?v=vGPI\\_JvLoN0](https://www.youtube.com/watch?v=vGPI_JvLoN0)).

Используем архитектуру *LeNet-5* и обучим сеть сначала на данных из набора *MNIST*.

In [1]:

```
! pip install tensorflow-gpu --pre --quiet  
! pip show tensorflow-gpu
```

Name: tensorflow-gpu  
Version: 2.2.0rc3  
Summary: TensorFlow is an open source machine learning framework for everyone.  
Home-page: <https://www.tensorflow.org/> (<https://www.tensorflow.org/>)  
Author: Google Inc.  
Author-email: packages@tensorflow.org  
License: Apache 2.0  
Location: /usr/local/lib/python3.6/dist-packages  
Requires: google-pasta, wrapt, wheel, protobuf, astunparse, h5py, gast, keras-preprocessing, six, grpcio, absl-py, opt-einsum, tensorboard, scipy, numpy, tensorflow-estimator, termcolor  
Required-by:

In [0]:

```
import tensorflow as tf  
from tensorflow import keras
```

In [0]:

```
import numpy as np
```

In [0]:

```
from tensorflow.keras.datasets import mnist  
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

In [0]:

```
x_train, x_test = tf.keras.utils.normalize(x_train, axis = 1), tf.keras.utils.normalize(x_t
```

In [0]:

```
x_train, x_test = x_train[..., np.newaxis], x_test[..., np.newaxis]
```

In [7]:

```
from tensorflow.keras.utils import to_categorical  
y_train, y_test = to_categorical(y_train), to_categorical(y_test)  
y_train.shape
```

Out[7]:

```
(60000, 10)
```

In [0]:

```
IMAGE_DIM_0, IMAGE_DIM_1 = x_train.shape[1], x_train.shape[2]
```

In [0]:

```
CLASSES_N = y_train.shape[1]
```

In [10]:

```
x_train.shape, x_test.shape
```

Out[10]:

```
((60000, 28, 28, 1), (10000, 28, 28, 1))
```

In [0]:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import AveragePooling2D, Conv2D, Dense, Flatten

model = tf.keras.Sequential()

model.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
                input_shape = (IMAGE_DIM_0, IMAGE_DIM_1, 1)))
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding =
model.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model.add(Flatten())
model.add(Dense(120, activation = 'tanh'))
model.add(Dense(84, activation = 'tanh'))
model.add(Dense(CLASSES_N, activation = 'softmax'))
```

In [0]:

```
# 'sparse_categorical_crossentropy' gave NAN loss

model.compile(optimizer = 'adam',
              loss = 'categorical_crossentropy',
              metrics = ['categorical_accuracy'])
```

In [13]:

```
model.summary()
```

Model: "sequential"

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| =====                        |                    |         |
| conv2d (Conv2D)              | (None, 28, 28, 6)  | 156     |
| -----                        |                    |         |
| average_pooling2d (AveragePo | (None, 14, 14, 6)  | 0       |
| -----                        |                    |         |
| conv2d_1 (Conv2D)            | (None, 10, 10, 16) | 2416    |
| -----                        |                    |         |
| average_pooling2d_1 (Average | (None, 5, 5, 16)   | 0       |
| -----                        |                    |         |
| flatten (Flatten)            | (None, 400)        | 0       |
| -----                        |                    |         |
| dense (Dense)                | (None, 120)        | 48120   |
| -----                        |                    |         |
| dense_1 (Dense)              | (None, 84)         | 10164   |
| -----                        |                    |         |
| dense_2 (Dense)              | (None, 10)         | 850     |
| =====                        |                    |         |
| Total params: 61,706         |                    |         |
| Trainable params: 61,706     |                    |         |
| Non-trainable params: 0      |                    |         |
| -----                        |                    |         |

In [0]:

```
EPOCHS_N = 20
```

In [15]:

```
model.fit(x = x_train, y = y_train, validation_split = 0.15, epochs = EPOCHS_N)
```

Epoch 1/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.2767 - categorical\_accuracy: 0.9161 - val\_loss: 0.1310 - val\_categorical\_accuracy: 0.9619

Epoch 2/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.1169 - categorical\_accuracy: 0.9639 - val\_loss: 0.0892 - val\_categorical\_accuracy: 0.9724

Epoch 3/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0789 - categorical\_accuracy: 0.9751 - val\_loss: 0.0922 - val\_categorical\_accuracy: 0.9732

Epoch 4/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0603 - categorical\_accuracy: 0.9808 - val\_loss: 0.0775 - val\_categorical\_accuracy: 0.9769

Epoch 5/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0465 - categorical\_accuracy: 0.9852 - val\_loss: 0.0781 - val\_categorical\_accuracy: 0.9759

Epoch 6/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0376 - categorical\_accuracy: 0.9877 - val\_loss: 0.0611 - val\_categorical\_accuracy: 0.9811

Epoch 7/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0305 - categorical\_accuracy: 0.9902 - val\_loss: 0.0683 - val\_categorical\_accuracy: 0.9793

Epoch 8/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0273 - categorical\_accuracy: 0.9911 - val\_loss: 0.0615 - val\_categorical\_accuracy: 0.9831

Epoch 9/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0224 - categorical\_accuracy: 0.9928 - val\_loss: 0.0674 - val\_categorical\_accuracy: 0.9812

Epoch 10/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0196 - categorical\_accuracy: 0.9936 - val\_loss: 0.0629 - val\_categorical\_accuracy: 0.9829

Epoch 11/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0165 - categorical\_accuracy: 0.9945 - val\_loss: 0.0683 - val\_categorical\_accuracy: 0.9826

Epoch 12/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0166 - categorical\_accuracy: 0.9942 - val\_loss: 0.0647 - val\_categorical\_accuracy: 0.9838

Epoch 13/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0121 - categorical\_accuracy: 0.9961 - val\_loss: 0.0799 - val\_categorical\_accuracy: 0.9794

Epoch 14/20

1594/1594 [=====] - 5s 3ms/step - loss: 0.0124 - categorical\_accuracy: 0.9957 - val\_loss: 0.0695 - val\_categorical\_accuracy: 0.9833

Epoch 15/20

```
1594/1594 [=====] - 5s 3ms/step - loss: 0.0112 - categorical_accuracy: 0.9964 - val_loss: 0.0743 - val_categorical_accuracy: 0.9821
Epoch 16/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0107 - categorical_accuracy: 0.9964 - val_loss: 0.0716 - val_categorical_accuracy: 0.9818
Epoch 17/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0102 - categorical_accuracy: 0.9964 - val_loss: 0.0757 - val_categorical_accuracy: 0.9838
Epoch 18/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0101 - categorical_accuracy: 0.9966 - val_loss: 0.0710 - val_categorical_accuracy: 0.9850
Epoch 19/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0080 - categorical_accuracy: 0.9973 - val_loss: 0.0765 - val_categorical_accuracy: 0.9846
Epoch 20/20
1594/1594 [=====] - 5s 3ms/step - loss: 0.0093 - categorical_accuracy: 0.9971 - val_loss: 0.0721 - val_categorical_accuracy: 0.9842
```

Out[15]:

```
<tensorflow.python.keras.callbacks.History at 0x7f0ad01780b8>
```

In [16]:

```
results = model.evaluate(x_test, y_test)

print('Test loss, test accuracy:', results)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.0756 - categorical_accuracy: 0.9830
Test loss, test accuracy: [0.0756305530667305, 0.9829999804496765]
```

Удалось достичь отличного результата — точность распознавания на тестовой выборке составила 98,0%.

## Задание 2

После уточнения модели на синтетических данных попробуйте обучить ее на реальных данных (набор *Google Street View*). Что изменилось в модели?

In [0]:

```
DS_URL_FOLDER = 'http://ufldl.stanford.edu/housenumbers/'

FIRST_DS_EXT = '.tar.gz'
SECOND_DS_EXT = '_32x32.mat'

TRAIN_DS_NAME = 'train'
TEST_DS_NAME = 'test'
EXTRA_DS_NAME = 'extra'
```

In [0]:

```
from urllib.request import urlretrieve
import tarfile
import os

def load_file(_url_folder, _name, _ext, _key, _local_ext = ''):

    file_url_ = _url_folder + _name + _ext

    local_file_name_ = _name + '_' + _key + _local_ext

    urlretrieve(file_url_, local_file_name_)

    return local_file_name_

def tar_gz_to_dir(_url_folder, _name, _ext, _key):

    local_file_name_ = load_file(_url_folder, _name, _ext, _key, _ext)

    dir_name_ = _name + '_' + _key

    with tarfile.open(local_file_name_, 'r:gz') as tar_:
        tar_.extractall(dir_name_)

    os.remove(local_file_name_)

    return dir_name_
```

In [0]:

```
second_ds_train_file = load_file(DS_URL_FOLDER, TRAIN_DS_NAME, SECOND_DS_EXT, 'second')
second_ds_test_file = load_file(DS_URL_FOLDER, TEST_DS_NAME, SECOND_DS_EXT, 'second')
second_ds_extra_file = load_file(DS_URL_FOLDER, EXTRA_DS_NAME, SECOND_DS_EXT, 'second')
```

In [0]:

```
from scipy import io

second_ds_train = io.loadmat(second_ds_train_file)
second_ds_test = io.loadmat(second_ds_test_file)
second_ds_extra = io.loadmat(second_ds_extra_file)
```

In [21]:

```
X_second_ds_train = np.moveaxis(second_ds_train['X'], -1, 0)
X_second_ds_test = np.moveaxis(second_ds_test['X'], -1, 0)
X_second_ds_extra = np.moveaxis(second_ds_extra['X'], -1, 0)

y_second_ds_train = second_ds_train['y']
y_second_ds_test = second_ds_test['y']
y_second_ds_extra = second_ds_extra['y']

print(X_second_ds_train.shape, y_second_ds_train.shape)
print(X_second_ds_test.shape, y_second_ds_test.shape)
print(X_second_ds_extra.shape, y_second_ds_extra.shape)
```


```
(73257, 32, 32, 3) (73257, 1)
(26032, 32, 32, 3) (26032, 1)
(531131, 32, 32, 3) (531131, 1)
```

In [22]:

```
import matplotlib.pyplot as plt

plt.imshow(X_second_ds_train[0])

plt.show()
```



In [0]:

```
IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2 = X_second_ds_train.shape[-3], X_second_ds_train
```

In [0]:

```
y_second_ds_train_cat = to_categorical(y_second_ds_train)
y_second_ds_test_cat = to_categorical(y_second_ds_test)
```

In [0]:

```
CLASSES_N_2 = y_second_ds_train_cat.shape[1]
```



In [0]:

```
model_2 = tf.keras.Sequential()

model_2.add(Conv2D(6, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2)))
model_2.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_2.add(Conv2D(16, kernel_size = (5, 5), strides = (1, 1), activation = 'tanh', padding
                  input_shape = (IMAGE_DIM_0_2, IMAGE_DIM_1_2, IMAGE_DIM_2_2)))
model_2.add(AveragePooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
model_2.add(Flatten())
model_2.add(Dense(120, activation = 'tanh'))
model_2.add(Dense(84, activation = 'tanh'))
model_2.add(Dense(CLASSES_N_2, activation = 'softmax'))
```

In [0]:

```
model_2.compile(optimizer = 'adam',
                loss = 'categorical_crossentropy',
                metrics = ['categorical_accuracy'])
```

In [28]:

```
model_2.summary()
```

Model: "sequential\_1"

| Layer (type)                 | Output Shape       | Param # |
|------------------------------|--------------------|---------|
| =====                        |                    |         |
| conv2d_2 (Conv2D)            | (None, 32, 32, 6)  | 456     |
| average_pooling2d_2 (Average | (None, 16, 16, 6)  | 0       |
| conv2d_3 (Conv2D)            | (None, 12, 12, 16) | 2416    |
| average_pooling2d_3 (Average | (None, 6, 6, 16)   | 0       |
| flatten_1 (Flatten)          | (None, 576)        | 0       |
| dense_3 (Dense)              | (None, 120)        | 69240   |
| dense_4 (Dense)              | (None, 84)         | 10164   |
| dense_5 (Dense)              | (None, 11)         | 935     |
| =====                        |                    |         |
| Total params: 83,211         |                    |         |
| Trainable params: 83,211     |                    |         |
| Non-trainable params: 0      |                    |         |
| =====                        |                    |         |

In [29]:

```
model_2.fit(x = X_second_ds_train, y = y_second_ds_train_cat, validation_split = 0.15, epochs=15)
```

Epoch 1/20

1946/1946 [=====] - 6s 3ms/step - loss: 1.0282 - categorical\_accuracy: 0.6717 - val\_loss: 0.7178 - val\_categorical\_accuracy: 0.7772

Epoch 2/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.6711 - categorical\_accuracy: 0.7900 - val\_loss: 0.6110 - val\_categorical\_accuracy: 0.8098

Epoch 3/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.5652 - categorical\_accuracy: 0.8228 - val\_loss: 0.6023 - val\_categorical\_accuracy: 0.8105

Epoch 4/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.5118 - categorical\_accuracy: 0.8394 - val\_loss: 0.6266 - val\_categorical\_accuracy: 0.8065

Epoch 5/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.4731 - categorical\_accuracy: 0.8516 - val\_loss: 0.5390 - val\_categorical\_accuracy: 0.8373

Epoch 6/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.4438 - categorical\_accuracy: 0.8621 - val\_loss: 0.4970 - val\_categorical\_accuracy: 0.8503

Epoch 7/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.4253 - categorical\_accuracy: 0.8668 - val\_loss: 0.4794 - val\_categorical\_accuracy: 0.8572

Epoch 8/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3957 - categorical\_accuracy: 0.8765 - val\_loss: 0.4964 - val\_categorical\_accuracy: 0.8514

Epoch 9/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3862 - categorical\_accuracy: 0.8797 - val\_loss: 0.4827 - val\_categorical\_accuracy: 0.8560

Epoch 10/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3683 - categorical\_accuracy: 0.8839 - val\_loss: 0.4800 - val\_categorical\_accuracy: 0.8530

Epoch 11/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3730 - categorical\_accuracy: 0.8805 - val\_loss: 0.4484 - val\_categorical\_accuracy: 0.8658

Epoch 12/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3508 - categorical\_accuracy: 0.8909 - val\_loss: 0.4820 - val\_categorical\_accuracy: 0.8544

Epoch 13/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3317 - categorical\_accuracy: 0.8961 - val\_loss: 0.4633 - val\_categorical\_accuracy: 0.8620

Epoch 14/20

1946/1946 [=====] - 6s 3ms/step - loss: 0.3272 - categorical\_accuracy: 0.8958 - val\_loss: 0.4577 - val\_categorical\_accuracy: 0.8642

Epoch 15/20

```
1946/1946 [=====] - 6s 3ms/step - loss: 0.3204 - categorical_accuracy: 0.8983 - val_loss: 0.4435 - val_categorical_accuracy: 0.8688
Epoch 16/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.3038 - categorical_accuracy: 0.9041 - val_loss: 0.4707 - val_categorical_accuracy: 0.8622
Epoch 17/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.2952 - categorical_accuracy: 0.9065 - val_loss: 0.4846 - val_categorical_accuracy: 0.8547
Epoch 18/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.2945 - categorical_accuracy: 0.9065 - val_loss: 0.4746 - val_categorical_accuracy: 0.8637
Epoch 19/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.2885 - categorical_accuracy: 0.9084 - val_loss: 0.4615 - val_categorical_accuracy: 0.8706
Epoch 20/20
1946/1946 [=====] - 6s 3ms/step - loss: 0.2745 - categorical_accuracy: 0.9137 - val_loss: 0.4993 - val_categorical_accuracy: 0.8527
```

Out[29]:

```
<tensorflow.python.keras.callbacks.History at 0x7f0ab60cc160>
```

In [30]:

```
results = model_2.evaluate(X_second_ds_test, y_second_ds_test_cat)
print('Test loss, test accuracy:', results)
```

```
814/814 [=====] - 2s 2ms/step - loss: 0.5624 - categorical_accuracy: 0.8352
Test loss, test accuracy: [0.5623723268508911, 0.8352412581443787]
```

Прежде всего, в модели изменилось то, что добавился ещё один класс — *не распознано*.

Эти данные более сложны для распознавания, что повлияло на результат — точность распознавания на тестовой выборке составила 77,4%.

Теперь реализуем распознавание первого датасета — реальных изображений с несколькими цифрами и границами. Для этого потребуется реализация алгоритма YOLO.

In [0]:

```
from imageio import imread
import pandas as pd

def image_to_array(_image):
    try:
        array_ = imread(_image)

        return True, array_
    except:
        return False, None

def dir_to_dataframe(_dir_path):

    data_ = []

    files_ = sorted(os.listdir(_dir_path))

    for f in files_:
        file_path_ = os.path.join(_dir_path, f)

        can_read_, im = image_to_array(file_path_)

        if can_read_:
            data_.append(im)

    dataframe_ = pd.DataFrame()

    dataframe_['data'] = np.array(data_)

    return dataframe_
```

In [0]:

```
first_ds_train_dir = tar_gz_to_dir(DS_URL_FOLDER, TRAIN_DS_NAME, FIRST_DS_EXT, 'first')
first_ds_test_dir = tar_gz_to_dir(DS_URL_FOLDER, TEST_DS_NAME, FIRST_DS_EXT, 'first')
```

In [0]:

```
first_ds_train_subdir = os.path.join(first_ds_train_dir, 'train')
first_ds_test_subdir = os.path.join(first_ds_test_dir, 'test')
```

In [0]:

```
first_ds_train_images_df = dir_to_dataframe(first_ds_train_subdir)
first_ds_test_images_df = dir_to_dataframe(first_ds_test_subdir)
```

In [0]:

```
import h5py

first_ds_train_boxes_mat = h5py.File(os.path.join(first_ds_train_subdir, 'digitStruct.mat'))
first_ds_test_boxes_mat = h5py.File(os.path.join(first_ds_test_subdir, 'digitStruct.mat'),
```

In [0]:

```
import numpy as np
import pickle
import h5py

def mat_to_pickle(_mat_path, _key):

    f = h5py.File(_mat_path, 'r')

    metadata = {}

    metadata['height'] = []
    metadata['label'] = []
    metadata['left'] = []
    metadata['top'] = []
    metadata['width'] = []

    def print_attrs(name, obj):
        vals = []
        if obj.shape[0] == 1:
            vals.append(int(obj[0][0]))
        else:
            for k in range(obj.shape[0]):
                vals.append(int(f[obj[k][0]][0][0]))
            metadata[name].append(vals)

    for item in f['/digitStruct/bbox']:
        f[item[0]].visititems(print_attrs)

    with open('{}pickle'.format(_key), 'wb') as pf:
        pickle.dump(metadata, pf, pickle.HIGHEST_PROTOCOL)
```

In [0]:


```
mat_to_pickle(os.path.join(first_ds_train_subdir, 'digitStruct.mat'), 'train_bbox')
mat_to_pickle(os.path.join(first_ds_test_subdir, 'digitStruct.mat'), 'test_bbox')
```

In [0]:

```
train_bbox_data = np.load('train_bbox.pickle', allow_pickle = True)
test_bbox_data = np.load('test_bbox.pickle', allow_pickle = True)
```

In [39]:

```
plt.imshow(first_ds_train_images_df['data'][0])  
plt.show()
```



In [40]:

```
train_bbox_data['label'][0]
```

Out[40]:

```
[1, 9]
```

### Задание 3

Сделайте множество снимков изображений номеров домов с помощью смартфона на ОС *Android*. Также можно использовать библиотеки *OpenCV*, *Simple CV* или *Pygame* для обработки изображений с общедоступных камер видеонаблюдения (например, <https://www.earthcam.com/>) (<https://www.earthcam.com/>).

В качестве примера использования библиотеки *TensorFlow* на смартфоне можете воспользоваться демонстрационным приложением от *Google* (<https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/android>) (<https://github.com/tensorflow/tensorflow/tree/master/tensorflow/examples/android>).

### Задание 4

Реализуйте приложение для ОС *Android*, которое может распознавать цифры в номерах домов, используя разработанный ранее классификатор. Какова доля правильных классификаций?